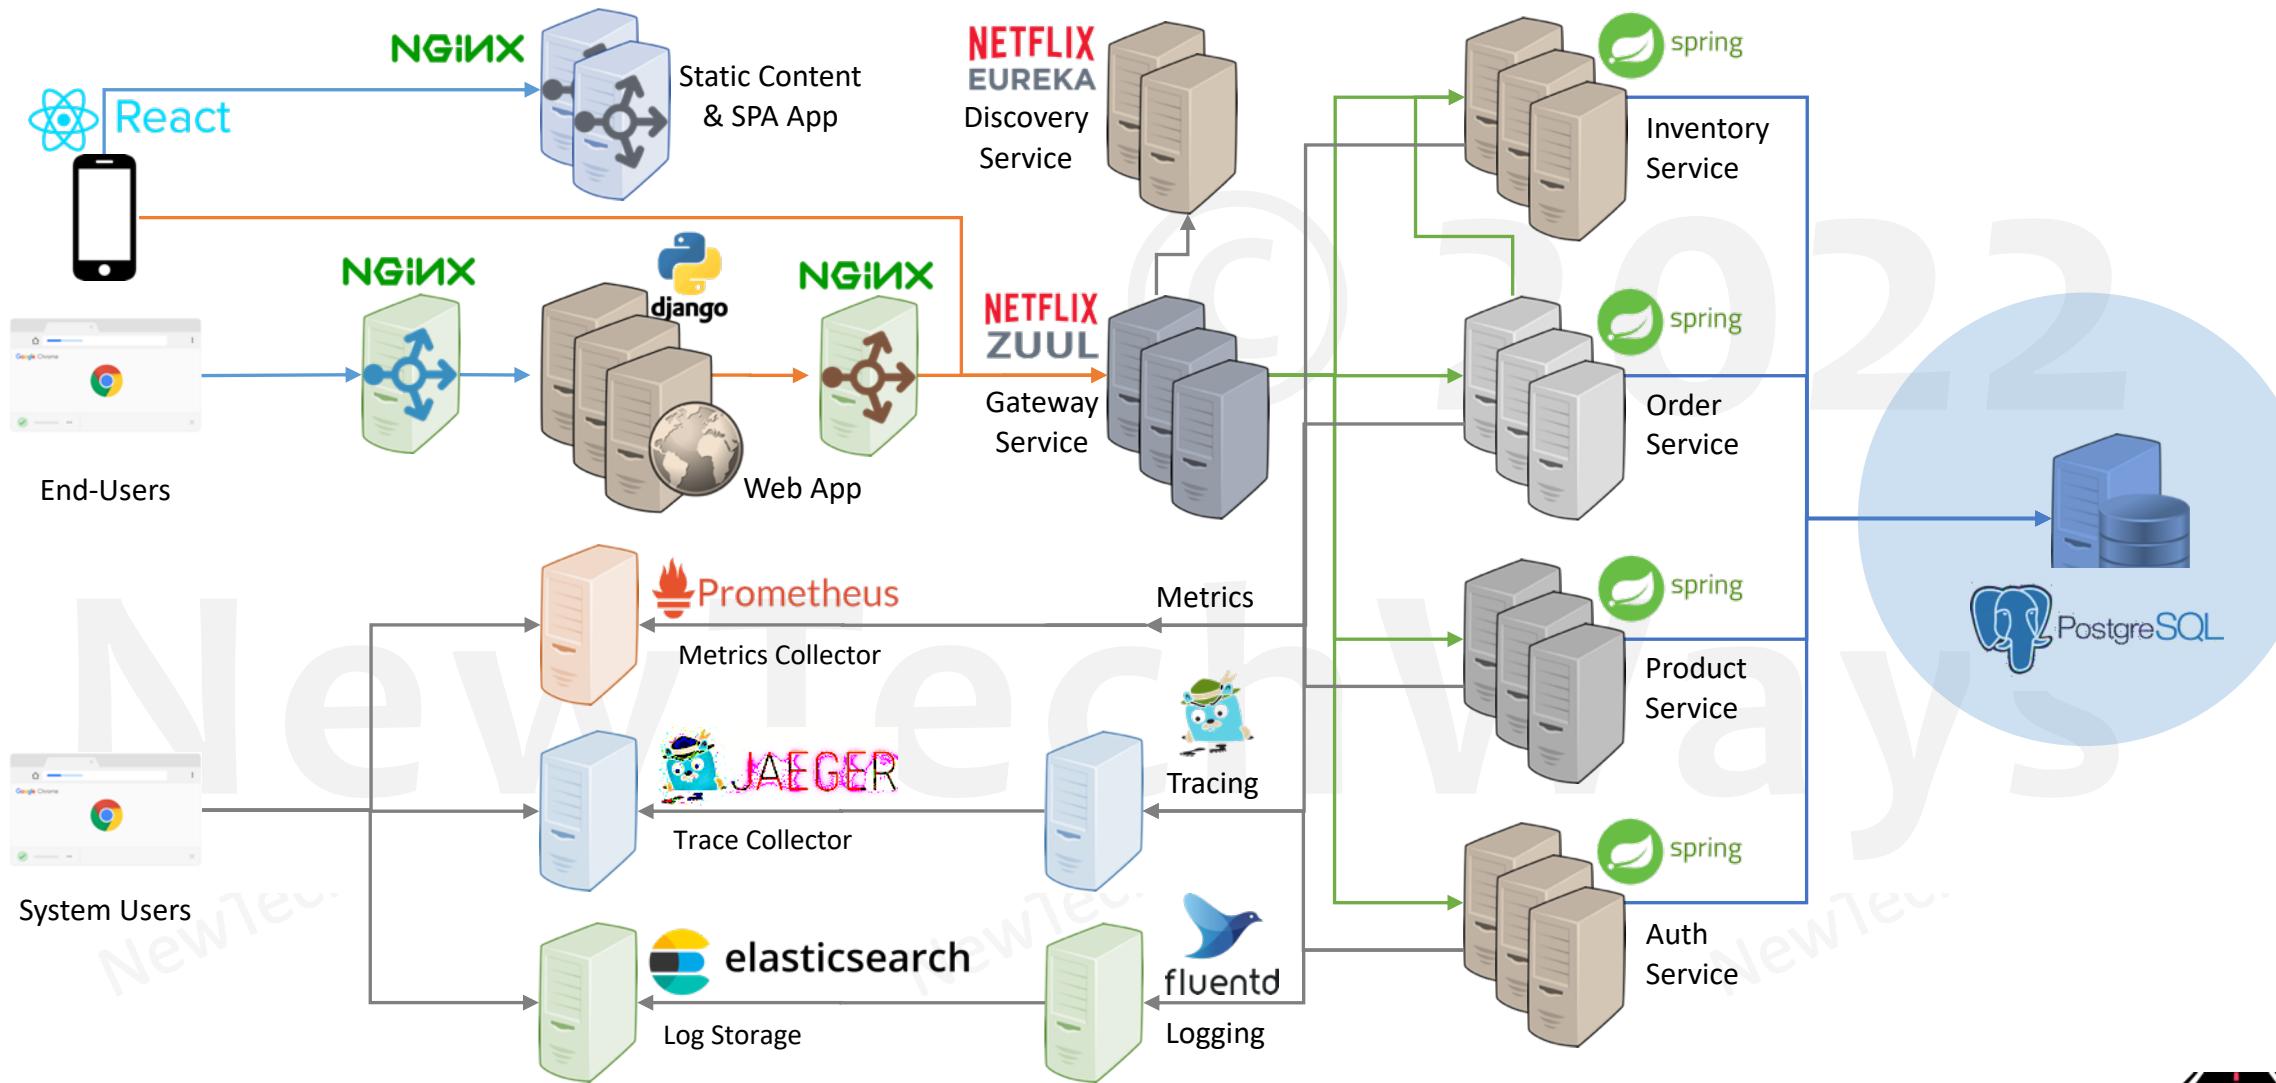


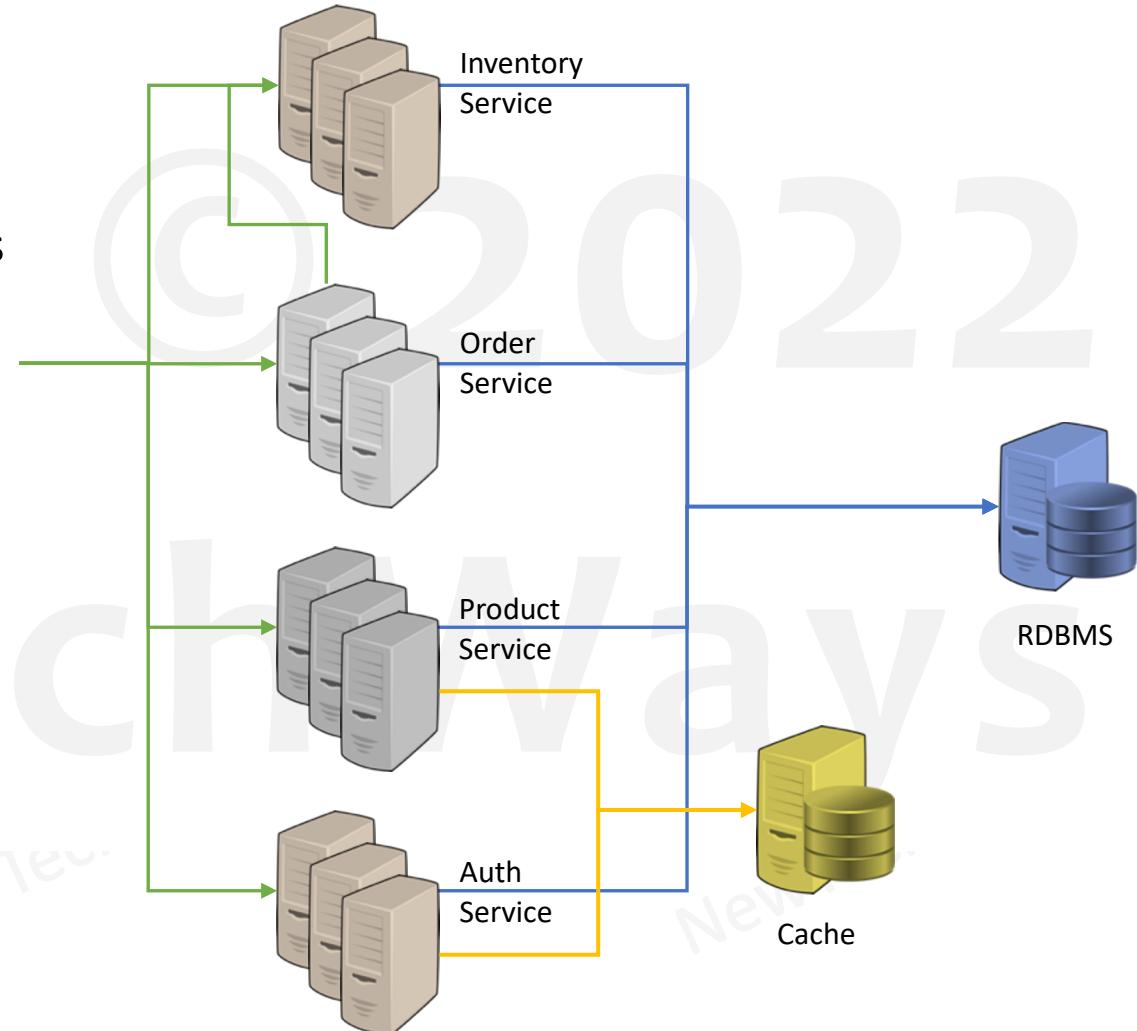
Data Management at High-Scale

High Data Load



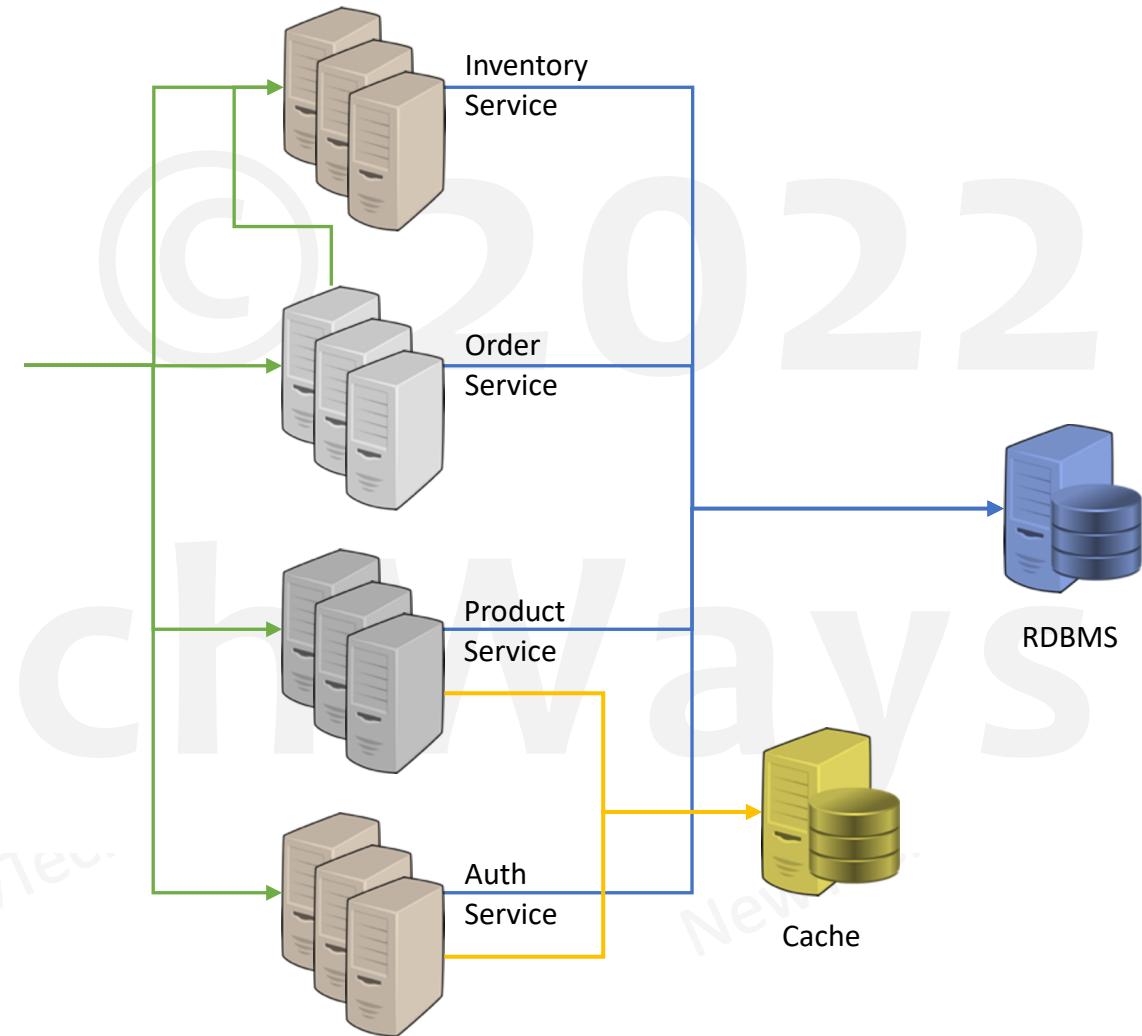
Managing Read Load – Caching

- Reduces data read load from the backend
 - Improve performance of read requests
 - Creates more capacity for write requests
- Challenge is to deal with stale data
 - TTL Values
 - Not ideal for dynamic data
 - Batch update
 - Read origin and update cache
 - Dynamic update
 - Update cache on add/update/delete



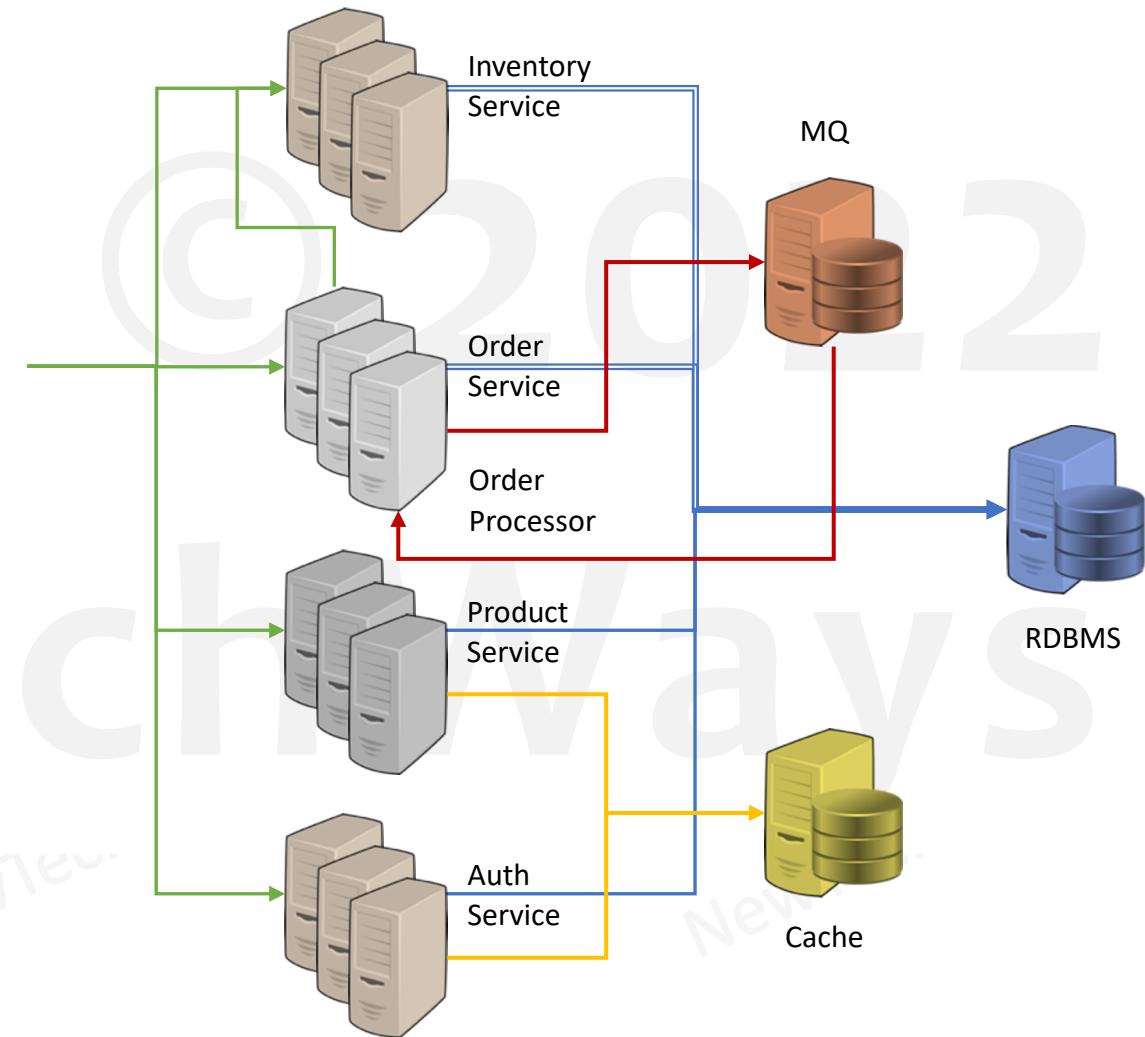
Managing Read Load – Caching

- Eviction Policy
 - LRU/LFU/TTL
 - Datasets that cannot fit in memory
 - View a product, user
 - No Eviction – Like DB
 - Datasets that can fit in memory
 - View all products in a category
- Cache Persistence
 - No frequent cache warmup
 - Restart – Maintenance/Failure



Managing Write Load – Asynchronous Processing

- Decouple write APIs from read APIs
 - Write APIs can be executed asynchronously
- Decouple order capture from order processing
 - Synchronous order capture
 - Low latency
 - Low failure rates
 - Improved user experience
 - Asynchronous order processing
 - Can wait for capacity availability
 - Asynchronous notification on completion

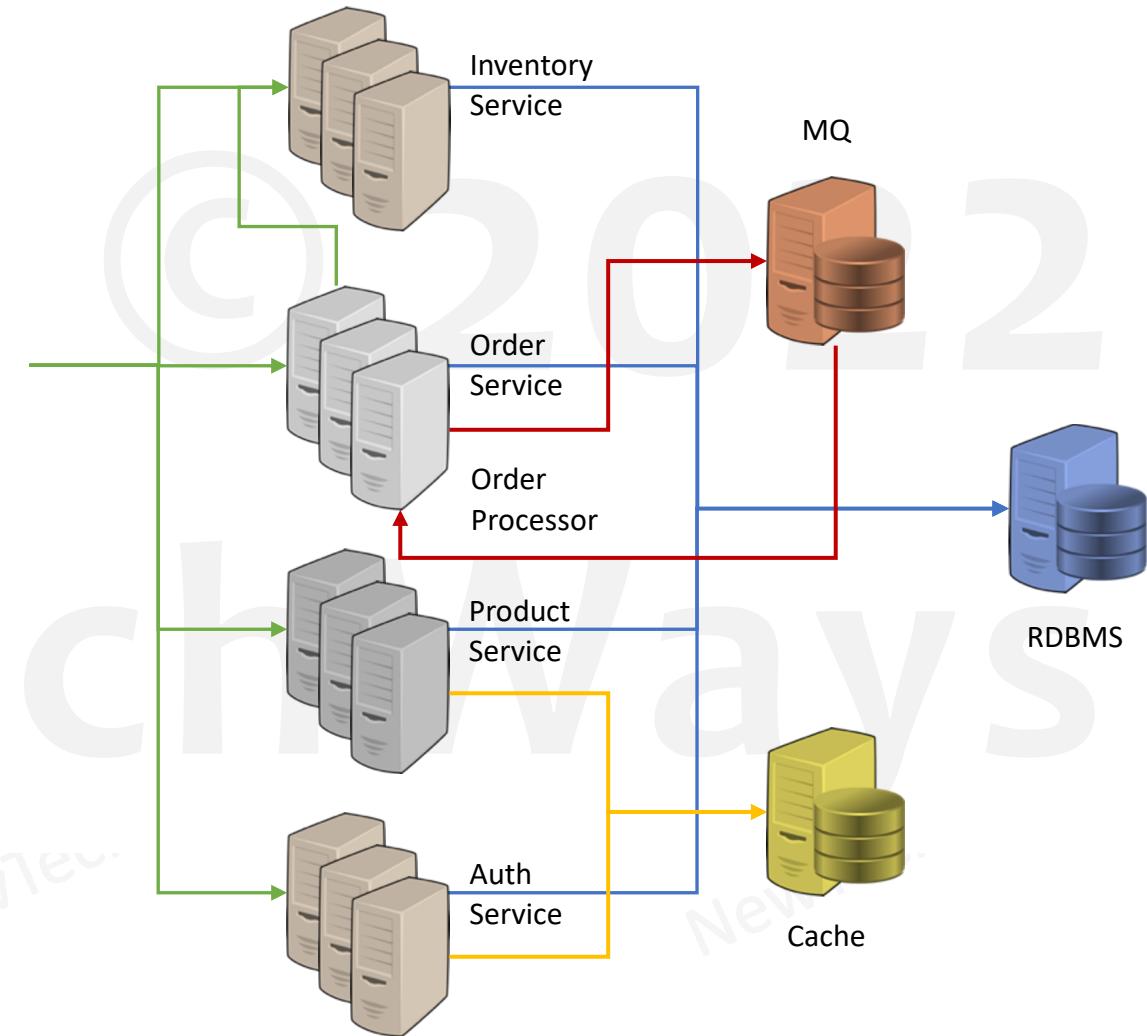


Managing Write Load – Message Queue

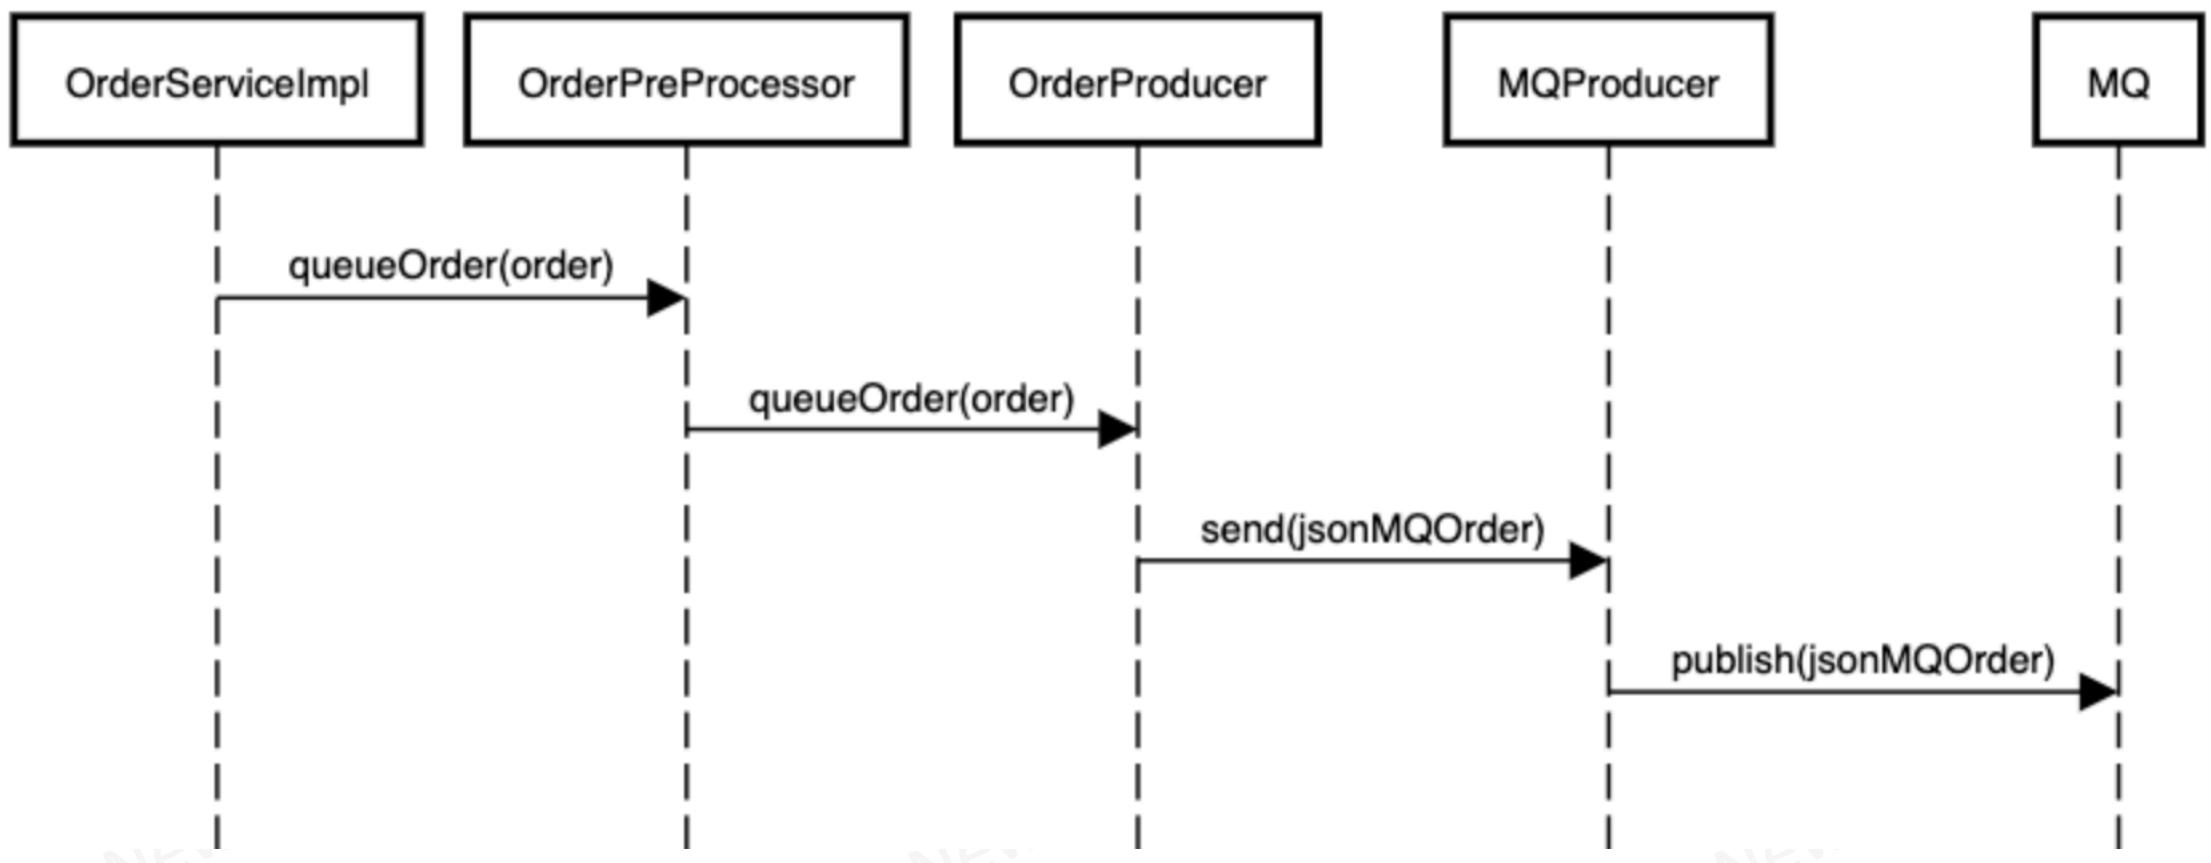
- Kafka is a distributed queue for high scalability
 - Partitioning of queue
- Kafka tradeoffs
 - Complex setup, requires Zookeeper
 - Message order only within a partition
 - No push message delivery
 - Pull requires constant polling



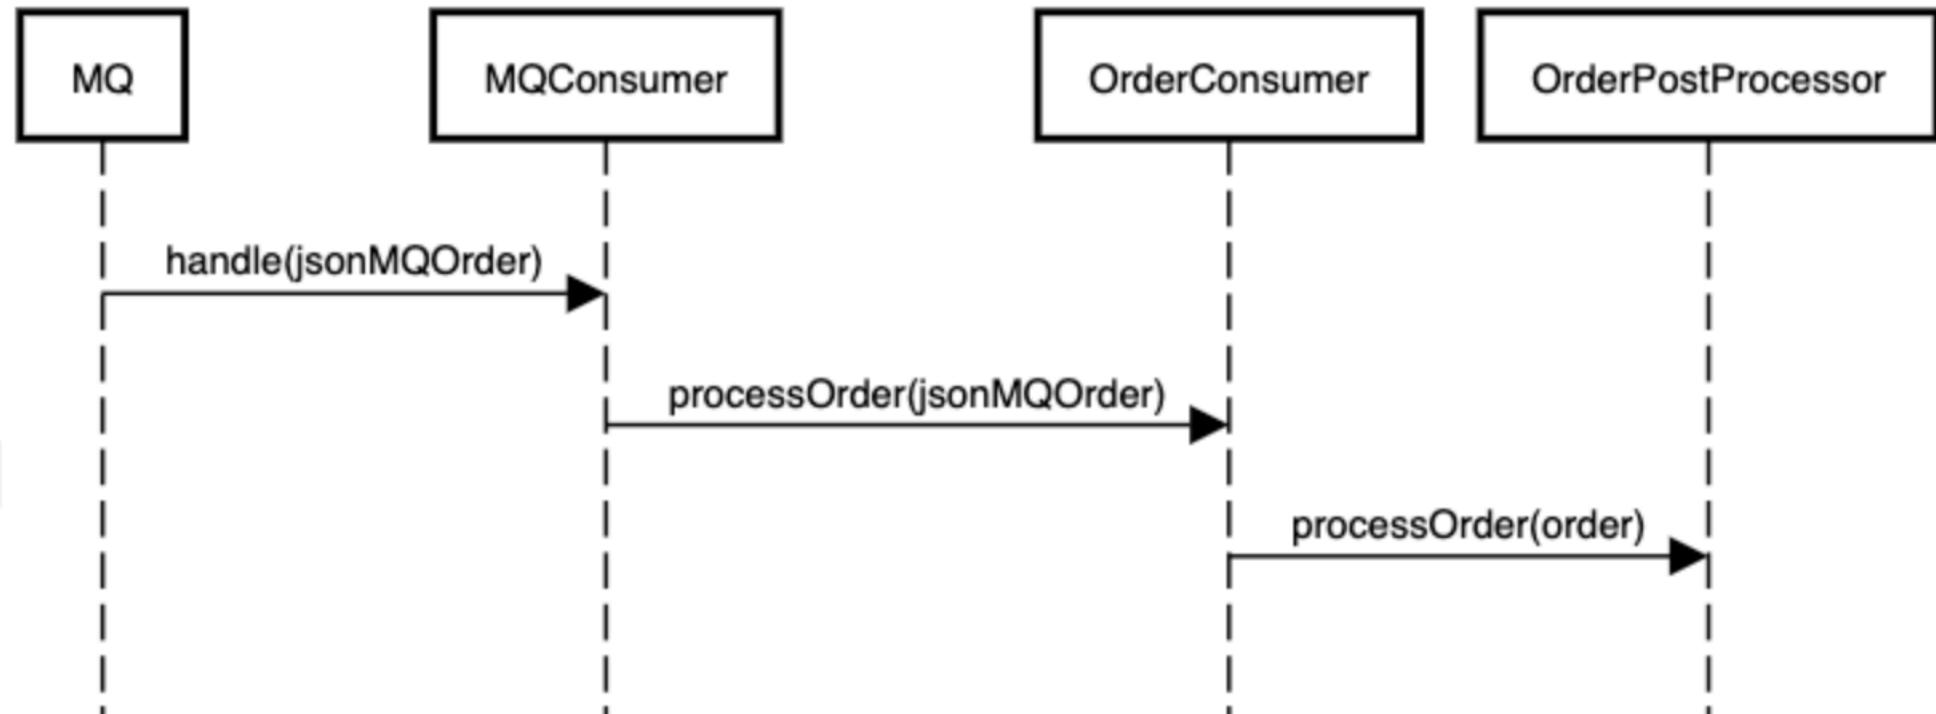
RabbitMQ™



Order Publish Flow

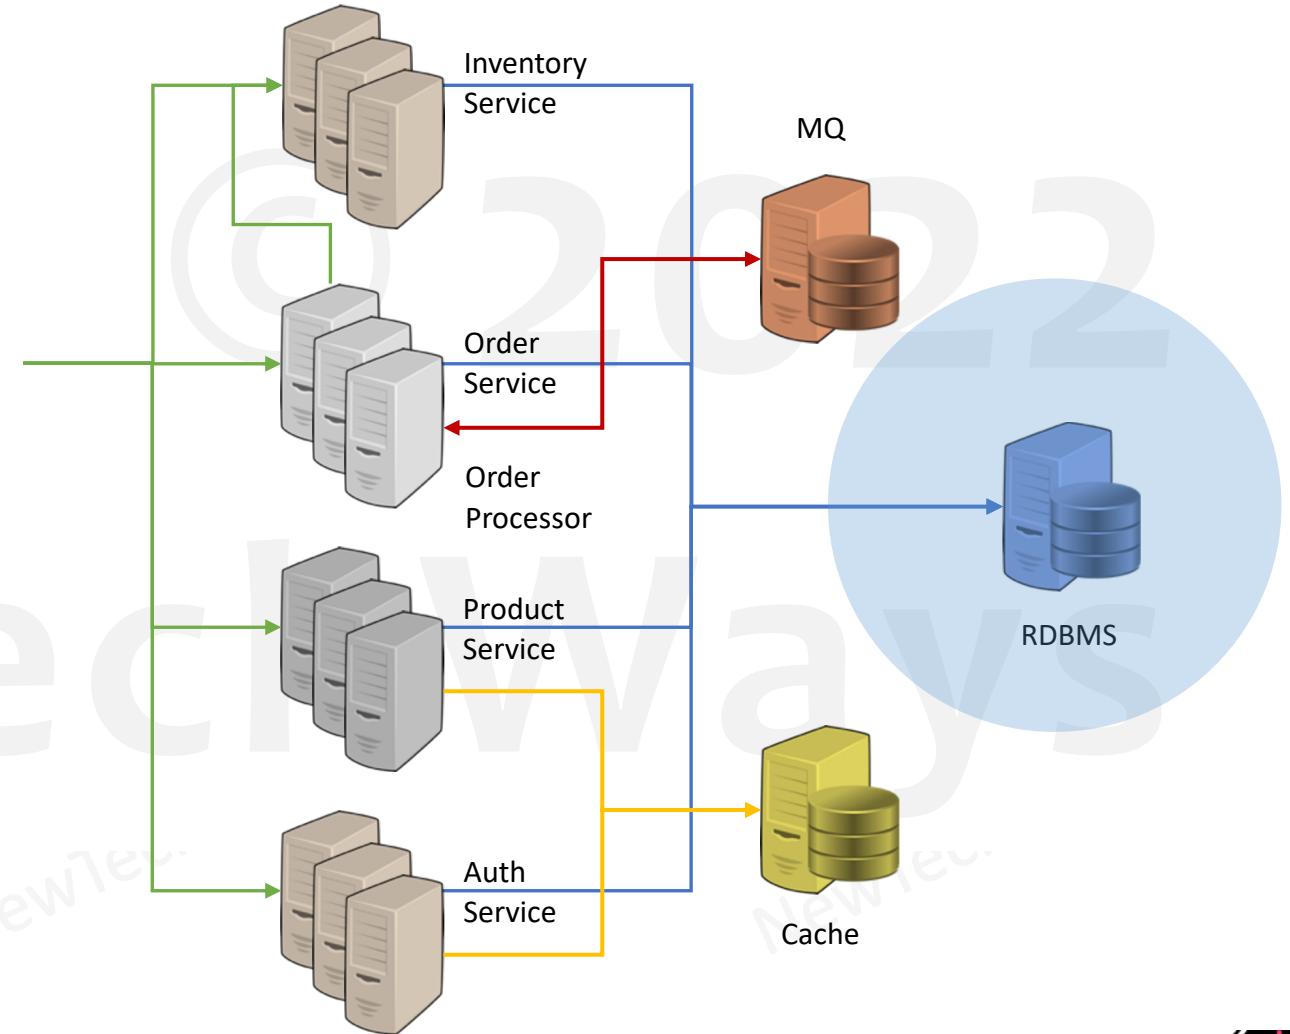


Order Consumer Flow



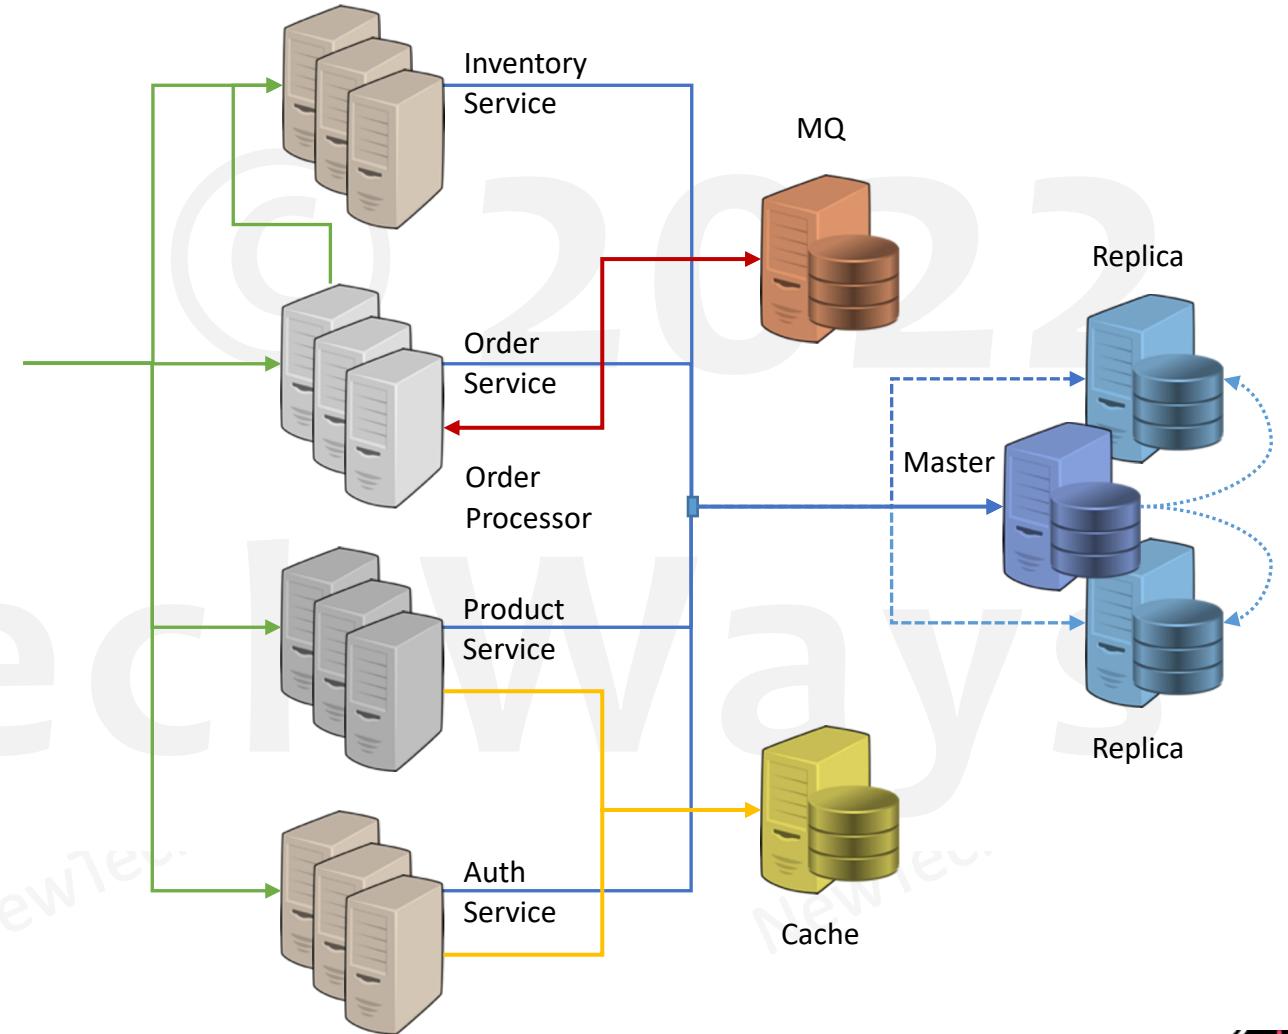
RDBMS Limitations

- RDBMS limitations
 - Single instance
 - < 1~5 TB of data



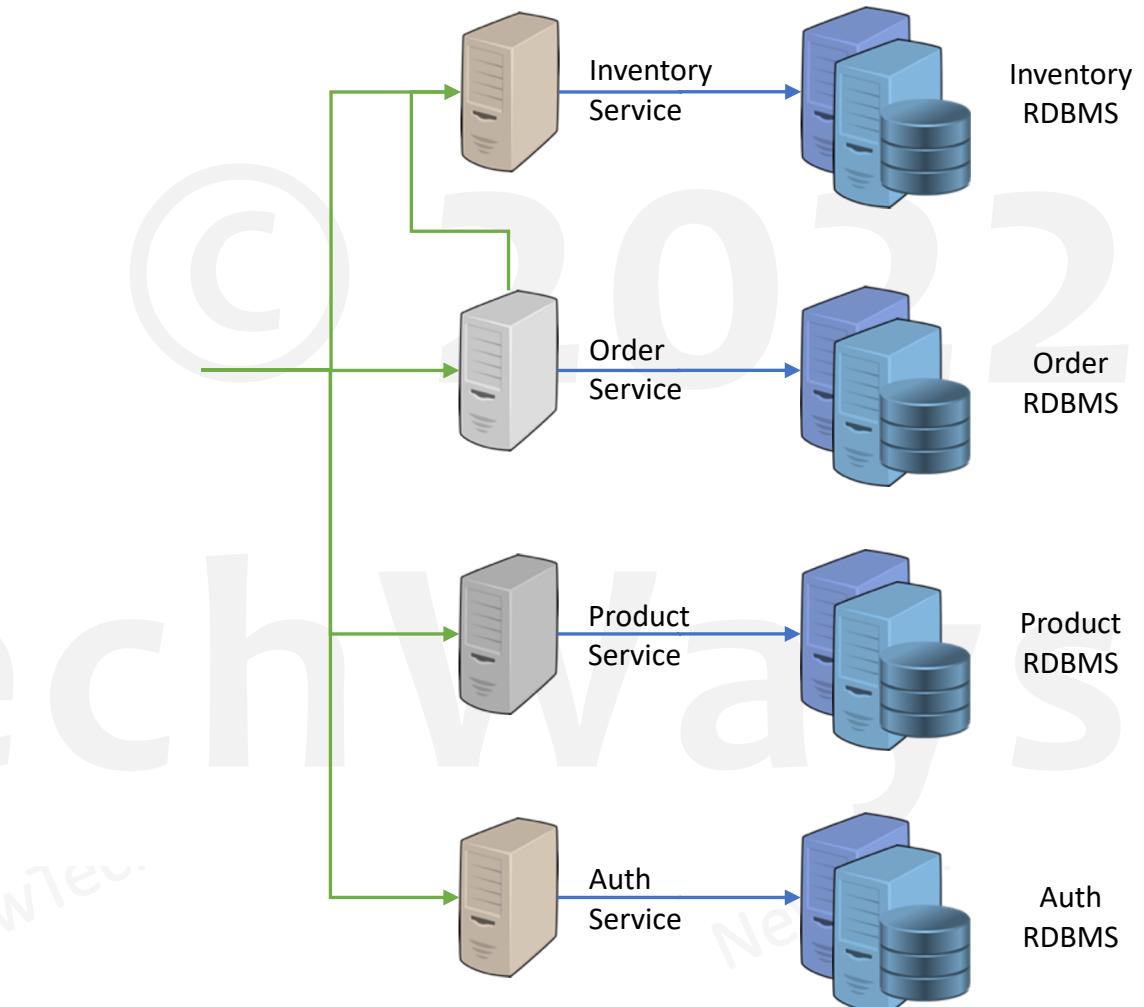
Partitioning Data Load

- Replication
 - High read load
 - Eventually consistent read queries
 - High availability
 - Write load is still a problem



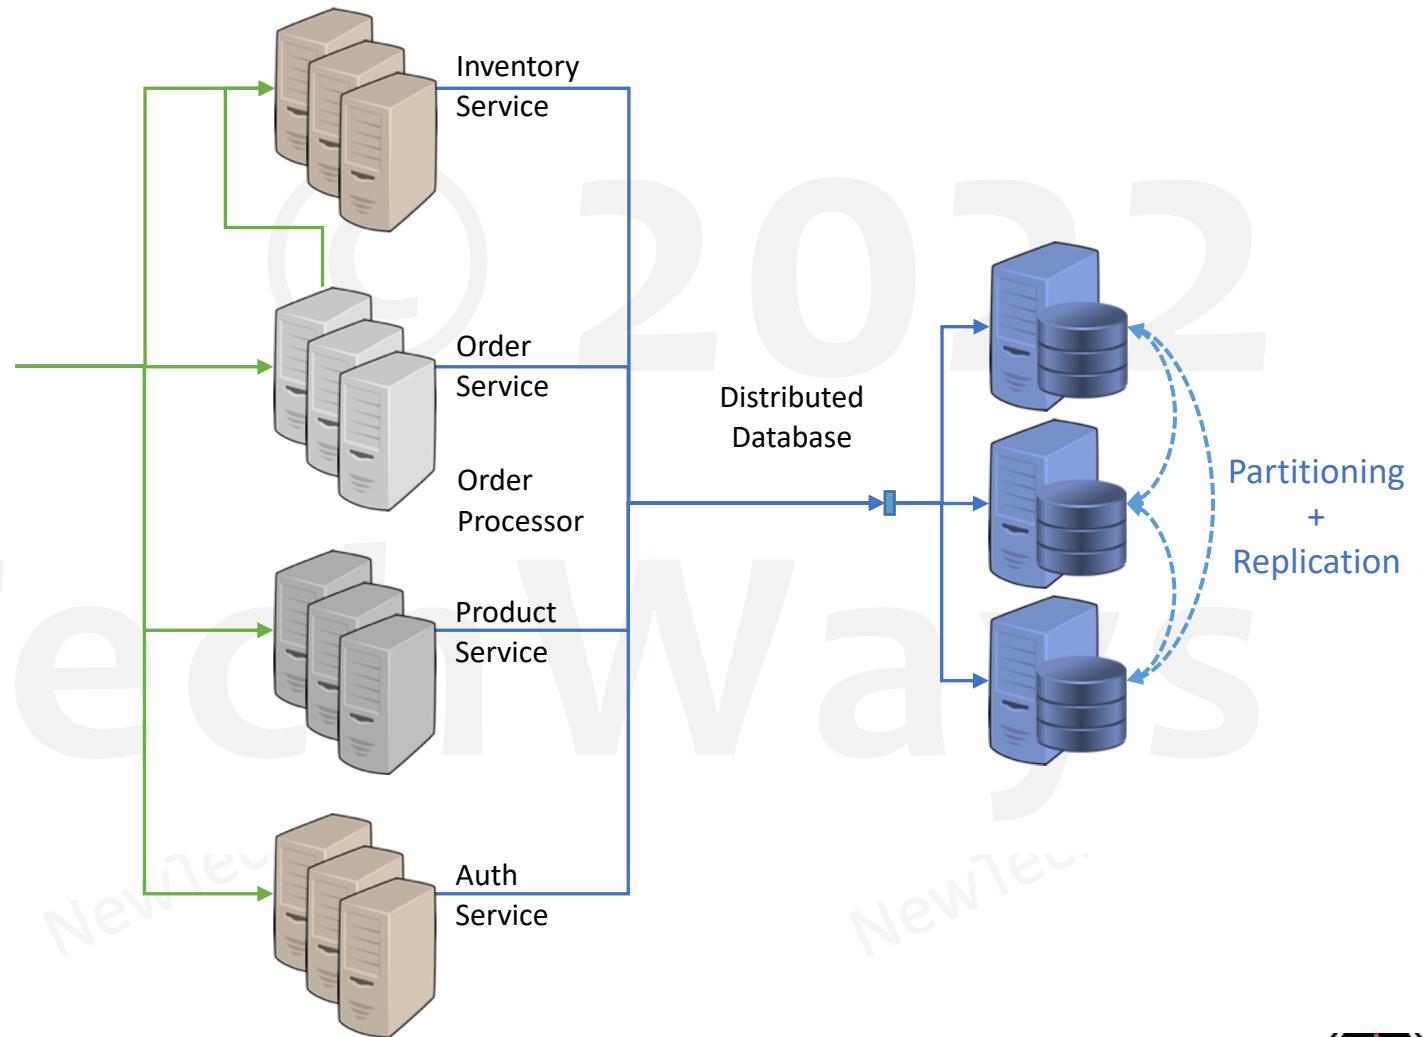
Partitioning – Application Level

- Vertical Partitioning
 - Microservices Architecture
 - Independent development and deployment
- Within a database
 - ACID, Joins, Constraints
- Across databases
 - No integrity constraints
 - No query joins
 - No ACID transactions
- Complexity of maintaining multiple databases



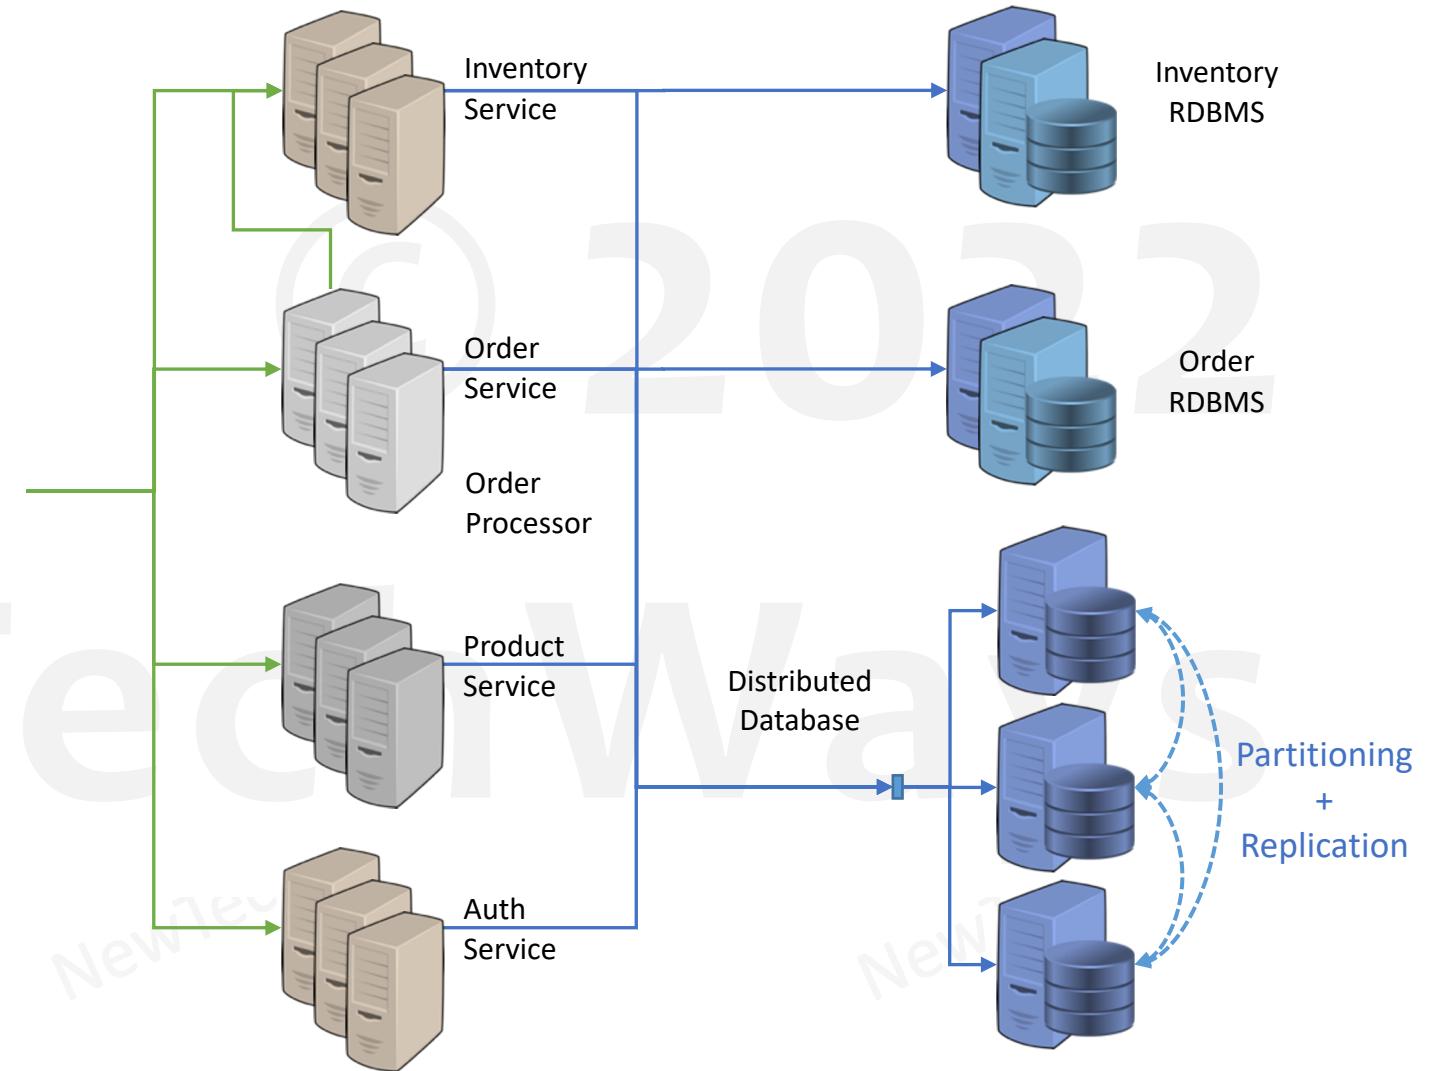
Partitioning – DB Record Level

- Data Partitioning
 - Record level
- Data Replication
- Routing request to serving instances
- Schema on demand
- Limited indexing
 - Puts limit on query patterns



Hybrid Approach

- RDBMS
 - For services that need strong consistency
- Distributed database
 - For services that can be eventually consistent



Partitioning Data Load

- MongoDB
 - Document Schema - JSON like
 - Hash Table like indexing
 - Master-Slave architecture
 - Single master for a record
- Cassandra
 - Column Family – Semi Structured Tables
 - Hash Table like indexing
 - Master-Master architecture - Write anywhere
 - Availability over consistency
- HBase
 - Column Family
 - B-Tree like indexing
 - Suitable for range queries
 - Master-Slave architecture



Partitioning Data Load

- MongoDB
 - Document type (b JSON) schema – Unstructured data
 - Hash Table like indexing
 - Master-Slave architecture
 - Write only on Master
- Cassandra
 - Table type schema – Structured data
 - SQL like CQL
 - Hash Table like indexing
 - Master-Master architecture
 - Write anywhere, no downtime
- HBase
 - Column Family type schema – Unstructured data
 - B-Tree like indexing
 - Suitable for analytics
 - Master-Slave architecture



© 2022



UserAuth	id	name	password	emailId
	<i>PK</i> (id)			

Cart	id
	<i>PK</i> (id)

UserRole	id	role
	<i>PK</i> (id, role)	

AUTH

Product	id	name	price
	<i>PK</i> (id)		

PRODUCT

Inventory	productId	Quantity
	<i>PK</i> (productId)	

INVENTORY

CartLine	cartId	cartLinId	productid	quantity
	<i>PK</i> (cartId, cartLinId)			

CART

Order	id	userId	date	status
	<i>PK</i> (id)			

ORDER

OrderLine	orderId	linId	productid	quantity
	<i>PK</i> (orderId, linId)			

UserAuth

id	name	password	emailId	role

*Partition Key (id)***AUTH****CartLine****CART**

cartId	cartLineId	productid	quantity

*Partition Key (cartId); Sort Key (cartLineId)***Product**

id	name	price

*Partition Key (id)***PRODUCT****Inventory****INVENTORY**

productId	Quantity

*Partition Key (productId)***OrderLine**

userId	id	orderId	lineld	productid	quantity	createdDate	createdTime	status

Partition Key (userId); Sort Key (orderId, lineld)