Interview Questions – Backup & Recovery Strategies

1. What is the difference between full, incremental, and differential backups? When would you use each?

• Full Backup:

- o Backs up all data, regardless of changes.
- o Pros: Easy to restore.
- o Cons: Time-consuming and requires high storage.
- Use case: Weekly full backups as a foundation.

• Incremental Backup:

- Backs up only the data that changed since the last backup (incremental or full).
- o Pros: Fast and storage-efficient.
- o Cons: Recovery requires all incremental backups.
- Use case: Daily backups after a full weekly backup.

• Differential Backup:

- o Backs up changes since the last full backup.
- o Pros: Faster restore than incremental.
- o Cons: Larger than incremental as the week progresses.
- **Use case**: Mid-week restore-friendly backups.

2. How do you define and balance RTO and RPO in a large-scale distributed system?

- RTO (Recovery Time Objective): Max acceptable time to restore service after failure.
- RPO (Recovery Point Objective): Max acceptable data loss measured in time (e.g., last 15 mins).

Balancing:

- Low RTO/RPO → Requires hot backups, replication, and higher cost.
- Analyze SLA commitments and criticality of components.
- Use tiered strategy:
 - Mission-critical: low RTO/RPO (e.g., failover DB, real-time replication).
 - Less-critical: higher RTO/RPO (e.g., batch systems, cold backups).

3. Explain cold, warm, and hot recovery strategies with examples.

Cold Recovery:

- No pre-configured resources; systems must be rebuilt from backups.
- Example: Backup stored on tape or cold cloud storage.
- Slowest, cheapest.

• Warm Recovery:

- Some components (e.g., data, configs) pre-provisioned, but app not running.
- <u>6</u> Example: Standby server with recent backups, manual DB restore.
- Moderate speed & cost.

Hot Recovery:

- Fully functional redundant system with real-time syncing.
- Fastest, most expensive.

✓ 4. How would you implement a backup strategy for a microservices-based application hosted in the cloud?

- 1. Identify critical services and data stores (DBs, object storage, configs).
- 2. Use cloud-native backup tools (e.g., AWS Backup, GCP snapshots).
- 3. Apply different backup frequencies per service:
 - Full weekly, daily incrementals.
 - Real-time replication for critical DBs.
- 4. Automate with IaC or CI/CD pipelines (e.g., Terraform, GitHub Actions).
- 5. Encrypt backups and store in **multi-region S3/Blob buckets**.
- 6. Regularly **test restoration** (chaos engineering / DR drills).

5. What trade-offs do you consider when designing a backup and recovery system for a high-availability service?

- Cost vs. recovery speed: Hot backups increase cost.
- Complexity vs. maintainability: Incremental backups are efficient but harder to restore.
- Storage vs. retention: Long retention increases storage needs.
- Compliance vs. agility: Regulatory backups might need immutability and longer archives.
- RTO/RPO: Must match business tolerance.

6. How does cloud storage simplify or complicate backup and recovery strategies?

Simplifies:

- Elastic, durable storage (e.g., S3, Azure Blob).
- Built-in snapshotting (e.g., EBS, RDS).
- Lifecycle management (automatic tiering, archival).
- Geo-redundancy support.

Complicates:

- Vendor lock-in risks.
- Costs can spiral with high frequency or long retention.
- Need for access management & encryption.
- Cross-region data compliance challenges.

7. What are some best practices for backup automation and testing in production systems?

- Automate backups with scheduled jobs or cloud-native tools.
- Use infrastructure as code to provision backup policies.
- Test restore procedures regularly (runbook + chaos testing).
- Monitor backup success/failures via alerts.
- Encrypt backups and verify data integrity.
- Follow 3-2-1 Rule:
 - o 3 copies
 - o 2 different media
 - o 1 offsite (e.g., cloud or DR site)

8. How would you handle backup for a database with terabytes of data and minimal allowed downtime?

- Use **point-in-time recovery (PITR)** if supported (e.g., MySQL binlogs, PostgreSQL WAL).
- Leverage incremental or log-based backup instead of full backups daily.
- Use replication (read-replica) to offload backups.
- Take **online snapshots** (e.g., EBS or managed RDS snapshots).
- Compress, encrypt, and store to cold + warm tiers.
- Use parallelism and dedicated backup windows to minimize performance impact.