

Interview Questions – Concurrency & Parallelism

◆ Conceptual Questions

1. What is the difference between concurrency and parallelism?

- **Concurrency** is the ability of a system to handle multiple tasks at once (by managing context switching and task progress), even if not literally executing them at the same time.
 - **Parallelism** means executing multiple tasks **simultaneously**, typically on multiple CPU cores.
 - **Example:**
 - Concurrency: One-core CPU interleaving execution of multiple tasks.
 - Parallelism: Multi-core CPU executing multiple tasks in parallel.
-

2. How do threads differ from processes?

- **Processes** are isolated execution environments with their own memory space.
 - **Threads** share the same memory space within a process but have their own stack.
 - Threads are **lighter and faster** to create compared to processes.
 - Communication between threads is easier (via shared memory), but this also introduces risks like race conditions.
-

3. What is a thread pool, and why is it preferred over creating new threads?

- A **thread pool** is a collection of reusable threads that are used to execute tasks.
- **Benefits:**
 - Reduces overhead of creating/destroying threads.
 - Prevents excessive thread creation (which can exhaust system resources).

- Improves response time and system throughput.
 - Common in web servers, background workers, async processing engines.
-

◆ Practical Scenarios

4. How would you design a web server to handle thousands of concurrent requests?

- Use **asynchronous non-blocking I/O** (e.g., `async/await` in .NET or event loops in Node.js).
 - Employ a **thread pool** for CPU-bound work.
 - Use **reverse proxies/load balancers** to distribute load.
 - Utilize **connection pooling**, **caching**, and **queue-based job processing** for scalability.
-

5. Describe how you would implement background job processing in a scalable system.

- Use a **message queue** (e.g., RabbitMQ, Kafka, Azure Queue).
 - Workers consume tasks from the queue and process them asynchronously.
 - Scale horizontally by increasing worker instances.
 - Ensure **idempotency** and **retry logic** for failure handling.
 - Monitor via dashboards and alerting tools.
-

6. How would you debug and resolve a deadlock in a multithreaded application?

- Use tools like **process dump analyzers**, **logging thread stacks**, or **Visual Studio Debugger**.
- Identify locks and resource acquisition order.
- Look for **circular wait conditions**.

- Fix by:
 - Always acquiring locks in a consistent global order.
 - Using **timeout-based locks**.
 - Minimizing lock scope (lock only what's necessary).
 - Switching to **concurrent collections** or lock-free data structures.
-

♦ Pitfall Awareness

7. What is a race condition, and how can you prevent it?

- A **race condition** occurs when multiple threads access and modify shared data concurrently without proper synchronization, leading to unpredictable results.
 - Prevention:
 - Use **locks/mutexes** (e.g., `lock` in C#).
 - Use **thread-safe collections**.
 - Leverage **atomic operations** (e.g., `Interlocked`).
 - Consider using **functional/stateless programming** where possible.
-

8. How do you ensure thread-safe operations in a shared-memory environment?

- Synchronize access to shared resources using:
 - `lock` or `Monitor` in C#
 - `ReaderWriterLockSlim` for read-heavy scenarios
- Use thread-safe APIs (e.g., `ConcurrentDictionary`)
- Prefer **immutable objects** and avoid shared state where possible.
- Leverage **Task Parallel Library (TPL)** and `async/await` to avoid manual thread management.

♦ Bonus (Advanced)

9. How does the event loop work in Node.js or similar environments?

- Node.js uses a **single-threaded event loop** that:
 - Executes tasks from the **call stack**.
 - Offloads blocking tasks (like I/O or timers) to the **libuv thread pool**.
 - Callback functions are queued in the **event queue** and picked up when the stack is empty.
- Enables **high concurrency** without multi-threading.

10. What's the difference between parallelism using threads vs. async I/O?

- **Threads** are used for **CPU-bound** operations. Parallelism here means executing tasks across multiple CPU cores.
- **Async I/O** is used for **I/O-bound** operations. It doesn't block threads and uses callbacks/promises/futures to resume execution when I/O completes.
- Combining both:
 - Async I/O for fast, scalable networking
 - Threads (or parallelism) for compute-heavy operations