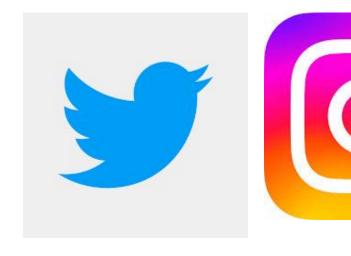
Mastering System Design

Design a News Feed (aka Twitter)

What is a News Feed?

- A timeline of updates/posts from people you follow.
- Core in platforms like Twitter, Instagram, Facebook.



Core Functional Requirements

- Post a tweet/update.
- Follow/unfollow users.
- View home timeline.
- Timeline must show updates from followed users.
- Likes, replies, retweets
- Media uploads (images, videos)
- Interaction Flow (User Perspective):
 - a. Alice follows Bob.
 - b. Bob tweets "Hello World".
 - c. Alice opens app \rightarrow sees Bob's tweet instantly.

Non-Functional Requirements

- Timeline expectations:
 - Feeds load instantly
 - New posts appear in real-time
 - Feed is relevant & fresh
 - Posts are not missing or duplicated
- High Availability
- Low Latency
- Scalability

Constraints & Challenges

- Fan-In / Fan-Out Challenge
 - 1 user \rightarrow 1 tweet \rightarrow must reach 1M followers?
 - Should we push to all followers? Or pull on demand?
- Pan-out on write vs. Fan-out on read this will shape our entire design.

Estimating Scale & Identifying Bottlenecks

- Daily Load Estimates:
 - o 👫 500M total users | 200M DAU

 - □ ~1B likes, 500M replies, 250M retweets
 - 2B+ feed requests/day (10 opens/user)

Read > Write: ~80% traffic is read-heavy

- Media Storage & Delivery
 - ~300M media uploads/day = 100–500TB/day
 - Must support:
 - Upload APIs
 - Object Storage (S3, GCS)
 - CDN-backed delivery
 - Tweet ↔ media linking

Identifying System Bottlenecks and Challenges

- Timeline = Fan-out at Scale
 - 1 user tweets → needs to be visible to 1M followers
 - Fan-out Models:
 - Fan-out on write: Pre-compute timelines (fast reads, heavy writes)
 - Fan-out on read: Compose timeline at read time (slower reads, lighter writes)
- 1 Hot users (celebrities) can trigger write storms
 - Read & Write Amplification
 - Write-Intensive Actions
 - Posting a tweet: May fan out to 100s or millions of timelines
 - Retweets, replies, likes: Update multiple engagement counters and visibility
 - Media uploads: Chunked writes, storage, and metadata association
 - Read-Heavy Patterns
 - Opening the timeline: Aggregates and sorts tweets from 100s of followed users
 - Viewing a tweet: Triggers fetches for media, replies, likes, retweets
 - Scrolling: Causes pagination, cache lookups, and lazy loads

One user action can trigger multiple backend reads/writes across services.

Core Services Breakdown

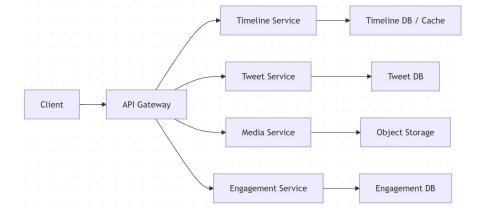
- User Service Profile, follow/unfollow
- **Tweet Service** Create, retrieve tweets
- **Timeline Service** Compose timelines for users
- **Engagement Service** Likes, replies, retweets
- Media Service Upload, store, fetch media
- Notification Service Fan-out alerts, activity
- Fanout Worker Background timeline propagation

High-Level Architecture (Zoomed Out)

User Flow:

 $\mathsf{Mobile/Web}\ \mathsf{App} \to \mathsf{API}\ \mathsf{Gateway} \to \mathsf{Microservices} \to \mathsf{DBs},\ \mathsf{Queues},\ \mathsf{Storage}$

- We use API Gateway for:
 - Request routing
 - Auth, rate limiting
 - Aggregating service calls



API Design (High level) - Sample API Contracts:

- POST /tweet
- GET /timeline
 - → returns latest tweets for a user
- POST /follow /unfollow
 - → modify following relationships
- POST /like /retweet /reply
 - → register engagement events
- POST /media/upload
 - → upload media, returns CDN URL

Note: All APIs are stateless; session/auth handled via token headers.

Timeline Generation Strategy

- Fan-out on Write
 - → Pre-compute timelines when a user tweets
 - Fast reads, heavy writes for hot users
- Fan-out on Read
 - → Compose timeline on demand from followed users
 - ✓ Lighter writes, ⚠ slower reads
- - Fan-out to regular users
 - Read-time fetch for hot users

Sync vs. Async Communication

- Synchronous (API Calls):
 - Fetch user timeline
 - Get tweet details
 - Submit engagement
- Asynchronous (Event Queues):
 - New tweet → enqueue fan-out jobs
 - Media uploads → process & link
 - Engagement → push notification trigger
- Improves throughput, decouples latency-sensitive paths

Strategic Tech & Infra Decisions

- @ Guiding Principles:
 - Optimize for read-heavy load, real-time delivery, and massive fan-out
 - Prioritize latency, scalability, and operational simplicity
- Key Tech Choices
 - Storage & DBs
 - Tweets, users, timelines → Scalable NoSQL (Cassandra, DynamoDB)
 - Engagements → Relational or Key-Value Store (PostgreSQL, Redis)
 - Media → Object Storage (Amazon S3, GCS) + CDN (Cloudflare, Akamai)
 - Messaging & Async Jobs
 - Kafka / RabbitMQ for fan-out jobs, engagement events, media pipelines
 - Caching
 - Redis / Memcached for hot timelines, tweets, and user sessions
 - o Compute & Infra
 - Kubernetes / ECS for autoscaling microservices
 - API Gateway / Envoy for routing, rate-limiting, auth
 - Monitoring & Resilience
 - Prometheus + Grafana for metrics, alerts
 - Circuit breakers, retries, queues for graceful degradation

📌 Every choice is made to handle massive volume, reduce latency, and isolate blast radius of failures.

The Final Design - News Feed (aka Twitter)

