
Mastering System Design

Design a Chat Application (aka WhatsApp)

What are we building? - And Functional Requirements







- We're designing a real-time chat application like WhatsApp.
- Functional Requirements
 - a. 1-to-1 and group messaging
 - b. Typing indicators and online status
 - c. Media support (images, videos, documents)
 - d. Delivery and read receipts
 - e. Multi-device support with sync



Non-Functional Requirements

- Sub-second message delivery
- High availability & fault tolerance
- End-to-end encryption
- Scalable to millions of users
- Smooth UX even on poor networks

Constraints & Challenges

-  Message Ordering & Delivery Guarantees
 - Ensure no duplicates, no lost messages, and correct order, even under retries or reconnections.
-  Unreliable Networks
 - Mobile users may frequently disconnect; we need robust retry, buffering, and sync mechanisms.
-  Security & Privacy
 - End-to-end encryption must be seamless, with no access to message content by servers.
-  Latency Expectations
 - Sub-second message delivery is expected, regardless of load or geography.
-  Presence & Sync
 - Track who's online and sync message state across all devices in near real-time.
-  High Scale
 - Millions of concurrent users, chatrooms, and messages per second; we must scale horizontally and efficiently.

Estimating Scale

- We begin with a few rough assumptions to get a sense of expected load:
 - Active users/day: 100 million
 - Average messages/user/day: 50
 - Messages/day: $100\text{M} \times 50 = 5$ billion messages/day
 - Peak traffic multiplier: 3× during global events
 - Concurrent connections: 20–30 million users online at peak

Identifying System Bottlenecks and Challenges

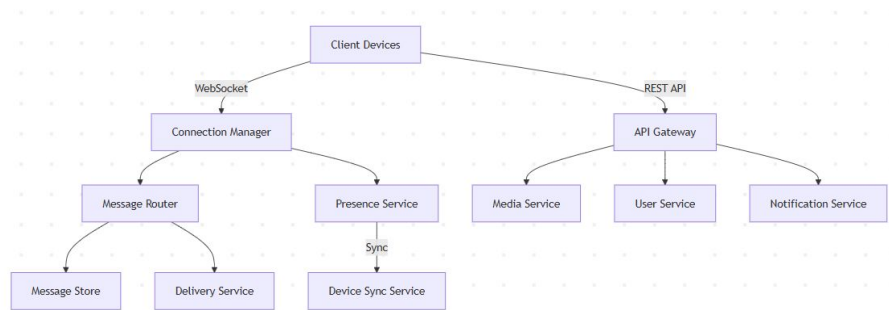
1. Message Ingestion & Fanout
 - a. Fanout to multiple recipients or devices causes exponential delivery load
 - b. Must queue, buffer, and batch smartly (e.g., Kafka, Redis Streams)
2. Presence Updates
 - a. Real-time online/offline tracking is chatty and high frequency
 - b. Use Pub/Sub patterns and efficient TTL-based caching
3. Delivery Acknowledgments
 - a. Each sent/read receipt adds more writes — can overload databases
 - b. Needs fast, write-optimized storage like Cassandra or DynamoDB
4. Sync Across Devices
 - a. Every message must sync across all logged-in devices
 - b. Requires a device registry and smart deduplication
5. Storage & Retrieval
 - a. Billions of messages stored with search capability
 - b. Must partition smartly and use append-friendly stores
6. Encryption & Security
 - a. End-to-end encryption adds CPU cost and complexity
 - b. Key exchange and secure metadata handling must scale too

Key Bottleneck: Real-Time Delivery (WebSockets)

- Real-time delivery is fundamental to chat systems — users expect messages to appear instantly. This means:
- Persistent connections need to be maintained for millions of users using WebSockets
- Load balancers and application servers must handle long-lived connections
- Need for connection management service: track connected users, devices, and routing
- Risk: connection churn, network drops, or mobile limitations (e.g., background app states)
- Why it's a bottleneck:
 - Handling millions of concurrent WebSocket connections requires optimized infra, horizontal scaling, and careful connection lifecycle management.

High-Level Architecture Overview and Key components

- **Connection Manager (WebSocket Service):**
 - Manages persistent WebSocket connections, delivers real-time messages, and tracks user/device sessions.
- **Chat Service:**
 - Core business logic for sending, receiving, storing messages, and managing delivery status and history.
- **Presence Service:**
 - Tracks online/offline status, typing indicators, and syncs across user devices using pub-sub or Redis.
- **Notification Service:**
 - Handles fallback notifications (e.g., push, SMS) when recipients are offline or WebSocket delivery fails.
- **Media Service:**
 - Manages upload and retrieval of media files like images, videos, and attachments using object storage.
- **Auth & User Service:**
 - Responsible for authentication, user profile management, and device binding.
- **Storage Layer:**
 - Backing stores for messages, presence, media metadata, and delivery receipts — optimized for high write throughput.
- **Group Service:**
 - Group Management: Handles group creation, deletion, and membership management (add/remove users).

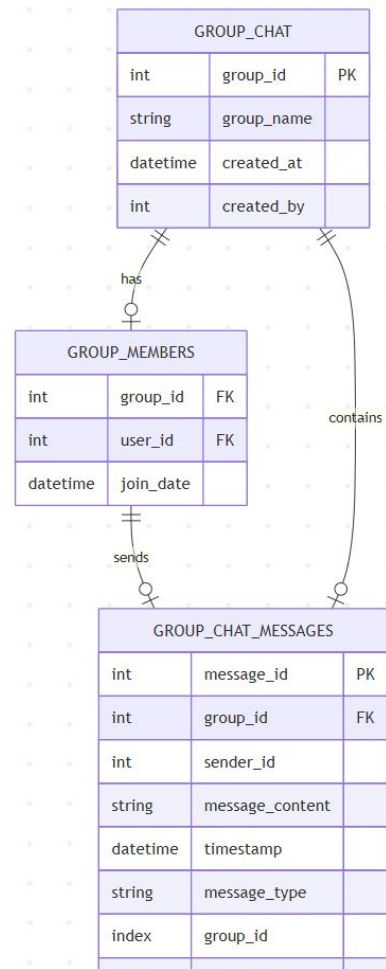


Real-Time Connection Manager (WebSocket Service)

- Maintains long-lived WebSocket connections per user and device.
- Handles message routing between senders and recipients.
- Integrates with Presence and Chat services for sync and state management.
- Supports scale-out using sticky sessions, Redis pub-sub, or message queues.
- Ensures delivery acknowledgments, retries, and device-level delivery tracking.

DB Schema for Direct Chat

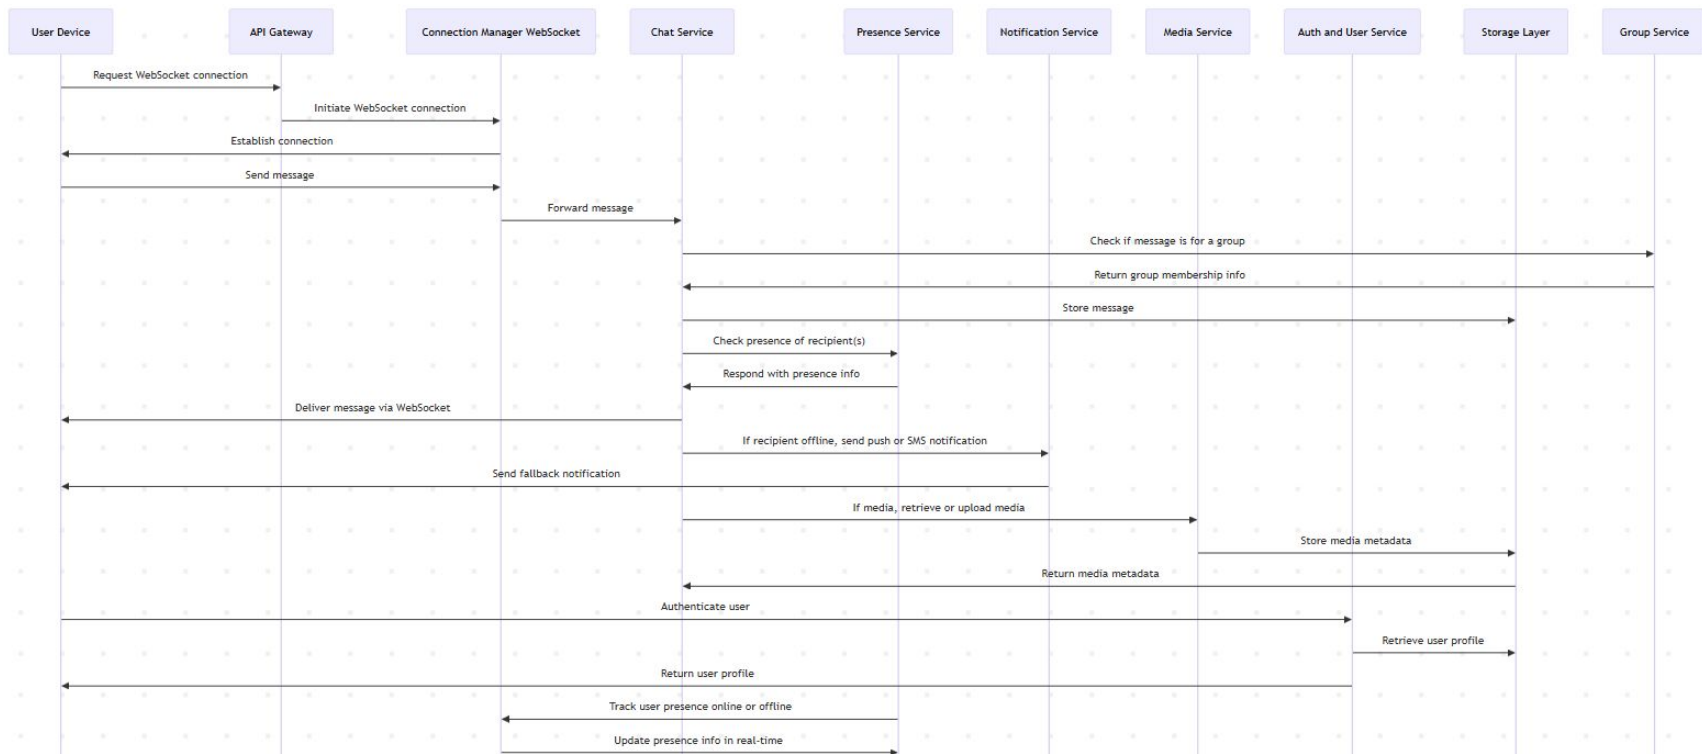
DIRECT_CHAT_MESSAGES		
int	message_id	PK
int	sender_id	
int	receiver_id	
string	message_content	
datetime	timestamp	
string	status	
string	message_type	
index	sender_id	
index	receiver_id	
index	timestamp	






How Chat & Group Chat Services Use Connection Manager for Real-Time Messaging

- **Direct Chat:**
 - Sender sends a message to the Chat Service with the recipient's user ID.
 - The Chat Service then sends the message to the Connection Manager.
 - Connection Manager looks up the recipient's active WebSocket connection and forwards the message to the corresponding WebSocket server.
 - The Receiver receives the message in real-time.
- **Group Chat:**
 - Sender sends a message to the Group Chat Service with the group ID and message content.
 - The Group Chat Service queries the Connection Manager, which looks up all active WebSocket connections for users in that group via the Group Members table.
 - The Connection Manager forwards the message to all group members' active WebSocket connections.
 - All group members receive the message simultaneously.

Chat Application - Message Flow & Sequence Diagram



API Design – WebSocket + REST

-  WebSocket Communication (Real-Time)
 - connect: Initiates persistent connection with access token and user metadata.
 - send_message: JSON payload with recipient ID, message content, timestamp, type.
 - message_ack: Acknowledgment from client upon delivery/read.
 - typing_indicator: Optional message indicating typing state.
-  REST APIs (Non-Real-Time Needs)
 - GET /messages?userId=&conversationId=
 - Fetch historical messages with pagination and filters.
 - POST /media/uploadUpload image/video/audio
 - attachments and get media URL.
 - GET /presence/{userId}
 - Fetch last seen / online status for a user.
 - POST /feedback
 - Submit crash reports or feedback (optional).
-  Security
 - All requests (WebSocket handshake + REST) require JWT auth.
 - Role-based access for user vs. admin endpoints.

Strategic Tech & Infra Decisions

- Tech Stack Choices:
 - Message Broker: Kafka (high-throughput) or AWS SQS (managed service)
 - Real-time Communication: WebSockets for instant messaging, leveraging SignalR (for .NET) or similar frameworks
 - Notification System: Firebase Cloud Messaging (FCM) for push notifications
 - Database: NoSQL databases (e.g., MongoDB) for fast reads and flexible schema
- Infrastructure & Deployment:
 - Deployment: Microservices on Kubernetes for scalability and management
 - Scaling: Horizontal scaling for WebSocket connections; KEDA (Kubernetes Event-Driven Autoscaling) for auto-scaling workers
 - Storage: Redis for caching active user sessions, PostgreSQL for transactional data (e.g., user settings)
- Security & Observability:
 - Authentication: OAuth2 for secure authentication across platforms
 - Logging & Monitoring: Prometheus and Grafana for real-time monitoring; CloudWatch for logs and alerts
 - High Availability: Multi-region deployment for failover; load balancing across WebSocket servers

The Final Design - Chat Application

