
Mastering System Design

Design an Online Rental Platform (aka Airbnb)

Understanding the Problem & Key Actors

- “Design an Online Rental Platform where property owners can list accommodations and users can search, book, and pay for short-term stays.”
- Core Concepts:
 - a. Listings (properties, availability, pricing)
 - b. Search and filters
 - c. Bookings and calendar synchronization
 - d. Payments and guest reviews
 - e. Media uploads (images, videos)
- Who Are the Key Actors?
 - a. Guest: Searches for and books accommodations
 - b. Host: Lists properties, manages availability and pricing
 - c. Admin: Moderates content, manages flagged users/listings
 - d. Payment Gateway: Processes secure transactions
 - e. External Calendar: Syncs host calendars (Google Calendar, iCal, etc.)



Functional Requirements

- Account creation & login (Host/Guest)
- Host can:
 - Create & update listings
 - Upload media
 - Set pricing and availability
- Guest can:
 - Search listings with filters
 - View listing details & reviews
 - Book listings and make payments
- Admins can manage users and moderate content
- Calendar sync with external platforms
- Notification system (email, push)

Non-Functional Requirements

- High Availability: 24/7 uptime, especially during peak travel seasons
- Scalability: Handle millions of users & listings
- Security: Payment security, personal data privacy
- Performance: <300ms response for key operations like search
- Reliability: Consistent booking logic, no double-booking
- Localization: Multiple currencies, time zones, and languages

Assumptions and Constraints

- Payments are handled via a 3rd-party gateway
- Reviews are moderated by the platform
- Users must verify email before booking/listing
- Media (images/videos) is uploaded to cloud object storage
- Users use web and mobile apps (we need REST APIs)
- Real-time search, but bookings can have slight delays (eventual consistency for availability)

Estimating Scale – Users, Listings, and Bookings

- Daily Active Users (DAU): ~5 million globally
- Concurrent Users: ~100k at peak
- Listings: 50 million+ (growing daily)
- Searches per day: 40–50 million
- Bookings per day: ~1 million
- Media per listing: Avg. 10 images, some videos
- Payment transactions/day: ~1 million (sync with bookings)

Implications:

- High volume of reads (search, view listings)
- Frequent writes (bookings, availability updates)
- Large media storage and delivery
- Real-time sync needed (calendar, availability)

Data Size & Storage Needs

- Rough Estimations:
 - Listings DB: $\sim 50\text{M listings} \times 5\text{ KB} = \sim 250\text{ GB}$
 - Bookings DB: $\sim 1\text{B records/year} \times 1\text{ KB} = \sim 1\text{ TB+}$
 - Media storage: $10\text{ images} \times 1\text{ MB} \times 50\text{M} = \sim 500\text{ TB}$
 - User profiles + history: $\sim 5\text{M DAU} \times 0.5\text{ KB} = \sim 2.5\text{ GB/day}$
- Hot Paths:
 - Searching listings (frequent reads)
 - Availability/calendar updates (write-heavy)
 - Booking confirmation (atomic write + payment + calendar update)
- Cold Paths:
 - Reviews history
 - User profile edits
 - Admin moderation

Identifying System Bottlenecks and Challenges

- Search Service
 - Handles very high query volume (QPS)
 - Needs fast, filtered, geo-based search
 - Requires scalable indexing and distributed querying
- Availability Calendar
 - Frequently updated due to bookings and sync with external calendars (iCal, Google Calendar, etc.)
 - Requires consistency to avoid double-bookings
 - Time zone management adds complexity
- Booking System
 - Requires atomic operation: availability lock + payment + confirmation
 - Needs to prevent race conditions during peak traffic
 - May benefit from queues or transactions to ensure reliability
- Media Storage & Delivery
 - Huge storage needs (images/videos for millions of listings)
 - High bandwidth and performance requirements for content delivery
 - Needs CDN caching and object storage (e.g., S3, Azure Blob)
- Payment Integration
 - High dependency on third-party APIs (Stripe, PayPal, etc.)
 - Must be fault-tolerant and secure (PCI compliance)
 - Needs retry logic, logging, and fallbacks

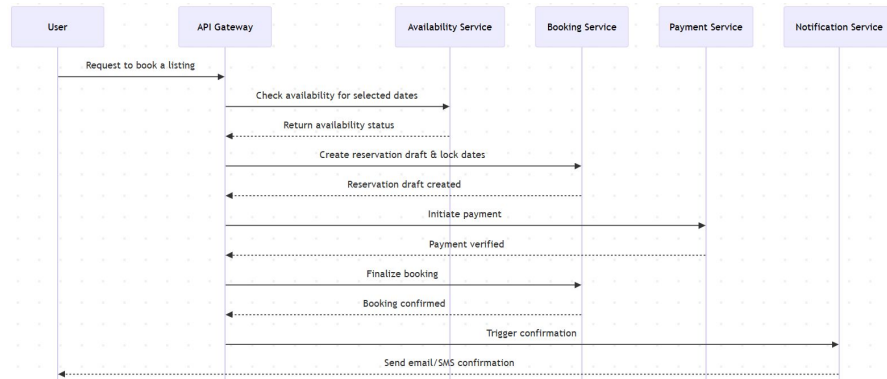
High-Level Architecture Overview and Key components

- **Frontend (Web & Mobile Apps):** User-facing clients for searching, booking, and managing listings
- **API Gateway:** Entry point for all client requests, Handles routing, authentication, rate limiting, and request shaping
- Core Microservices
 - **User Service:** Handles sign-up, login, profiles, preferences
 - **Listing Service:** Manages property listings (details, images, amenities)
 - **Search Service:** Indexes listings for geo-based, filtered search
 - **Availability Service:** Manages booking calendar and date availability
 - **Booking Service:** Handles reservation lifecycle and booking flow
 - **Payment Service:** Integrates with payment providers, tracks payments
 - **Notification Service:** Email, SMS, and push notifications
 - **Review & Ratings Service:** User feedback, property ratings
- Support Services
 - **Media Service:** Upload & store images/videos (via S3, CDN)
 - **Calendar Sync Service:** Syncs with external calendar providers
 - **Analytics & Logging Service:** Tracks user events and system health

Service Interactions (Example: Booking Flow)

- Step 1: User requests to book a listing
- Step 2: Availability Service checks if dates are open
- Step 3: Booking Service locks availability, creates a reservation draft
- Step 4: Payment Service initiates and verifies payment
- Step 5: Upon success, Booking Service finalizes booking
- Step 6: Notification Service sends confirmation email/SMS

All interactions are asynchronous where possible (e.g., payment confirmation, notifications) using message queues (e.g., RabbitMQ, Kafka).



Communication Patterns & APIs

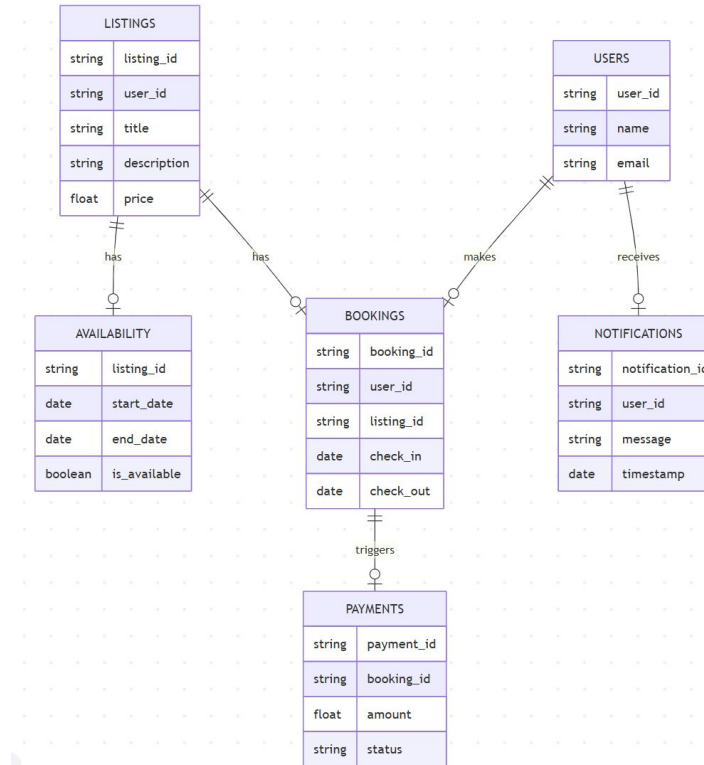
- Sync Communication (REST/gRPC):
 - Used for user interactions, search, listing fetch
 - APIs exposed via API Gateway
- Async Communication (Message Queue/Event Bus):
 - Booking events (e.g., booking_created, booking_failed)
 - Notifications, calendar syncs, email dispatches
 - Payments webhook events
- Authentication & Authorization
 - OAuth2/JWT-based tokens
 - Role-based access: guest, host, admin

Data Storage & Indexing Strategy

- Primary Datastores
 - Relational DB (PostgreSQL/MySQL) for transactions, user & listing data
 - NoSQL (MongoDB/DynamoDB) for availability snapshots, reviews
- Search Index
 - Elasticsearch for full-text + geo search on listings
- Caching Layers
 - Redis for hot data (recent searches, popular listings)
 - CDN (Cloudflare/Akamai) for images and static content

Sample DB Schema for Online Rental Platform

- Users
- Listings
- Availability
- Bookings
- Payments
- Notifications
- Media



Strategic Tech & Infra Decisions

- Scalability
 - Use Azure for flexible scaling and horizontal scaling across multiple servers.
- Redundancy & High Availability
 - Leverage load balancers and multi-region database replication for redundancy and availability.
- Data Storage
 - Store transactional data in SQL and media assets in Azure Blob Storage for durability and cost-effectiveness.
- Asynchronous Communication
 - Use RabbitMQ or Kafka for asynchronous tasks like payments and notifications to decouple services.
- Caching
 - Integrate Redis for caching frequently accessed data to improve response time and reduce load.
- Microservices Architecture
 - Adopt microservices to scale and maintain services (e.g., booking, payment, availability) independently.
- Monitoring & Logging
 - Use Azure Monitor for health tracking and centralized logging (ElasticSearch) for real-time diagnostics.

The Final Design - Online Rental Platform

