
Mastering System Design

Design a URL Shortener (aka TinyURL)

Introduction to TinyURL

- TinyURL is a URL shortening service that converts long URLs into short, unique links for easy sharing and tracking.
- Example:
 - `https://www.example.com/very/long/path/to/resource` → `https://tinyurl.com/abc123`
- Why it's Useful:
 - Compact Links: Perfect for platforms like Twitter and SMS
 - Improved UX: Easy to share, clean, and readable
 - Click Tracking: Analytics for link performance
 - Branding: Custom domains for enterprises
 - Cross-Channel Friendly: Works across emails, chats, and print
- How it Works:
 - User submits a long URL
 - System generates a unique short key
 - Stores mapping in database
 - Redirects short URL to original













Functional Requirements for TinyURL

- **Shorten a URL:** Accept a valid long URL and return a shortened URL.
 - Example: Input: `https://example.com/article?id=1234` → Output: `https://tinyurl.com/xYz12`
- **Redirect to Original URL:** When accessing the short URL, redirect to the original long URL.
- **Prevent Duplicate Short URLs:** If the same long URL is submitted, return the same short URL or handle according to configuration (unless custom alias is used).
- **User Authentication:** Allow users to register/login to manage URLs, view analytics, and set expiration.





Non-Functional Requirements for TinyURL

- **High Availability:** System must be available 24/7 with >99.9% uptime.
- **Performance & Low Latency:** URL redirection should occur in milliseconds. Shortening URLs should be near-instantaneous.
- **Scalability:** System must handle millions or billions of URLs, supporting high read volume (redirects) and moderate write volume (URL shortening).
- **Reliability:** Ensure data persistence and no data loss even during failures using durable storage and backups.

Unique URL Generation Strategies

- Random String Generation: Creates a fixed-length string from random characters
 -  Unpredictable, no obvious pattern
 -  Risk of collisions, requires collision handling
 -  Okay for unpredictability, but adds complexity
- UUID (Universally Unique Identifier): 128-bit globally unique identifier (e.g., 123e4567-e89b-12d3-a456...)
 -  Guaranteed uniqueness, no central coordination
 -  Very long, not user-friendly
 -  Not ideal for TinyURL due to length
- Hashing with Salt: Hashes the original URL (e.g., SHA-256 + salt)
 -  Unique, secure, hard to reverse
 -  May not be short, collision possible, needs mapping storage
 -  Useful for security-focused cases, but not optimal for shortening
- Base62 Encoding: Converts incrementing ID to Base62 (0-9, a-z, A-Z)
 -  Short, compact, deterministic, easy to implement
 -  Needs counter management to avoid collisions
 -  **Recommended for TinyURL (fast, scalable)**

Estimating Scale & Identifying Bottlenecks

-  **Estimated User Traffic**
 - Daily Active Users (DAU): ~10 million
 - Monthly Active Users (MAU): ~300 million
 - New Short URLs/day: ~100,000 (1% of DAU)
 - Redirect Requests/day: ~50 million (5 per user)
-  **Memory Requirement (Hot URL Cache)**
 - Cache top 1M most accessed URLs
 - Each mapping \approx 500 bytes
 - Total memory \approx 500 MB
-  **Network Bandwidth (Redirects)**
 - 50M redirects/day \times 700 bytes = ~35 GB/day
 - Avg throughput: ~0.4 MB/sec
 - Peak throughput: ~5 MB/sec
-  **Storage Requirement (URL Mapping DB)**
 - 100K new URLs/day \times 500 bytes = ~50 MB/day
 - Yearly data \approx 18 GB raw + overhead
 - Round up to ~50 GB/year (with indexes, logs, backups)

Bottleneck Identification

- High read volume → Focus on cache and fast DB reads
- Write throughput is moderate → Ensure consistency
- Latency sensitivity in redirects → Low-latency infra needed
- Plan for burst traffic with autoscaling & CDN support

API Design – Create Short URL

- Accepts a long URL and returns a shortened URL.
- Endpoint: **POST /api/shorten**
- Request (JSON):

```
{  
  "long_url": "https://www.example.com/article?id=123",  
}
```

- Response (JSON):

```
{  
  "short_url": "https://tinyurl.com/my-alias"  
}
```


API Design – Redirect

- Redirect to Original URL
- Endpoint: **GET /:short_key**
 - Example: GET /abc123
- Behavior:
 - Looks up the original long URL using the short key.
 - Returns a 302 HTTP Redirect to the long URL.

API Design – Delete

- Delete a Short URL (Optional)
- Endpoint: **DELETE /api/url/:short_key**
- Behavior:
 - Deletes the mapping if user is authenticated and owns the URL.
 - Requires Bearer token for authorization.

User Authentication APIs

- User Registration
- Endpoint: **POST /api/auth/register**

```
Request:
{
  "name": "John Doe",
  "email": "user@example.com",
  "password": "securePassword123"
}
Response:
{
  "message": "Registration successful"
}
```

- User Login
- Endpoint: **POST /api/auth/login**

```
Request:
{
  "email": "user@example.com",
  "password": "securePassword123"
}
Response:
{
  "access_token": "<JWT>",
  "token_type": "Bearer",
  "expires_in": 3600
}
```





- Secure Endpoints with Bearer Token:

```
Authorization: Bearer <access_token>
```

High-Level System Design – Overview

- API Gateway: Entry point for all clients; handles routing, rate limiting, auth
- URL Shortener Service: Contains logic for key generation, duplicate checking, alias validation
- Redirect Service: High-performance resolver for short keys → long URLs
- Database: Persistent store for all mappings, users, metadata
- Cache Layer: Redis/Memcached for top N frequently accessed URLs
- Auth Service: Manages user login, JWT tokens, and sessions

Collision Handling in Distributed URL Generation - Hello Zookeeper

- Why Collisions Happen
 - Multiple service instances generating IDs independently → risk of duplicates
 - No global coordination → Base62 encoding same ID → incorrect URL mapping
- What is Zookeeper?
 - Distributed coordination service by Apache
 - Ensures synchronization across nodes in a distributed system
- Zookeeper as a Solution
 -  Atomic ID generation using Zookeeper-managed global counter
 -  Guarantees each instance gets a unique ID
 -  Uses znodes to store and manage counters
 -  Supports distributed locking to serialize ID generation
- Flow
 - Service requests next ID from Zookeeper
 - Zookeeper increments global counter atomically
 - ID is Base62 encoded and used as TinyURL
 - (Optional) Mapping stored in DB

Strategic Tech & Infra Decisions for TinyURL

- Database:
 - SQL (e.g., PostgreSQL with auto-increment IDs)
 - NoSQL (e.g., Redis for caching, DynamoDB for scalability)
 - Cache: Redis or Memcached for high-speed lookup
- Scalability & Performance
 - Horizontal scaling for URL generation services
- High Availability
 - Load Balancer to distribute traffic across service instances
 - Replication in DB to avoid single points of failure
 - Failover-ready infrastructure using cloud-managed DBs or services

The Final Design - TinyURL

- URL Generation Flow
 - User submits long URL via API
 - API Gateway handles authentication, throttling
 - URL Generation Service:
 - Requests a unique ID from Zookeeper (ensures no collisions)
 - Encodes it (e.g., Base62) to create the short URL
 - Stores the mapping in the Database (Short URL → Long URL)
 - Response sent back with the generated TinyURL
- Redirection Flow
 - User hits a short URL (e.g., tinyurl.com/abc123)
 - API Gateway forwards to Redirection Service
 - Redirection Service:
 - Checks Cache for the short URL (fast-path)
 - If not found, queries the Database (cold-path)
 - Redirects user to the original long URL (HTTP 302)

