
Mastering System Design

Design an Auction Platform (aka eBay)

What is an Auction Platform?



- A digital system that enables users to list items and bid in real-time within a defined time window.
- Auctions promote competitive pricing and urgency — the highest bid at closing time wins.
- The platform handles item listings, bid tracking, real-time updates, auction lifecycle, and payment processing.
- Key Objective:
 - a. Deliver a secure, fair, and scalable real-time bidding experience — from listing to payment.








Functional Requirements

- User registration & authentication
- Item listing with auction parameters (start/end time, reserve price)
- Real-time bid placement and updates
- Auction state transitions: scheduled → active → ended
- Payment processing after auction ends (e.g., via Stripe, PayPal)
- Notifications: outbid, auction won, payment pending/complete






Key Actors & Use Cases

-  Actors:
 - Seller: Lists items, sets pricing rules
 - Bidder: Places bids, gets real-time notifications
 - System: Handles auction logic, enforces timing, manages payments
 - Admin: Monitors platform activity and handles fraud or disputes
-  Use Cases:
 - Seller lists a gaming console for a 3-day auction
 - Multiple bidders join and bid competitively
 - System tracks bids, updates UI in real-time
 - Winning bidder completes payment → seller is notified






Non-Functional Requirements

-  Performance:
 - Sub-second latency for bid placement and bid updates
-  Scalability:
 - Support thousands of concurrent auctions and real-time users
-  Security:
 - Secure user authentication, payment data protection, anti-bot protections
-  Availability:
 - Ensure high availability during peak auction times and payment flow
-  Observability:
 - Logging of auction and payment events, real-time monitoring dashboards




Constraints & Challenges

-  **Real-Time Pressure:**
 - Handle bids submitted within milliseconds of closing — resolve ties & order events accurately
-  **Concurrency Conflicts:**
 - Simultaneous bid updates must be conflict-free and idempotent
-  **Live Updates:**
 - Efficient, scalable real-time delivery to all watchers of an auction (via WebSockets/pub-sub)
-  **Payments Workflow:**
 - Payment failures, timeouts, retries, and fraud detection must be handled gracefully
-  **Fairness & Trust:**
 - Prevent bots/sniping and ensure fair auction rules are enforced transparently

Estimating Scale

- Assumed Metrics:
 -  Users: 5M registered users, 500K DAU
 -  Active Listings: 1M ongoing auctions at any moment
 -  Bids: Avg. 10 bids/auction → 10M bids/day
 -  Peak Activity: 10K concurrent users bidding/viewing one popular auction
 -  Payments: 100K completed auctions/day → 100K payment transactions/day

Traffic Patterns & Real-Time Pressure Points










-  Read-Heavy Operations:
 - Viewing auction listings & item details
 - Real-time bid updates to watchers
 - Browsing user profiles, search, filters
 - Scale: Millions of reads per minute at peak
-  Write-Sensitive Operations:
 - Bid placement (needs low latency, strong consistency)
 - Creating new auction listings
 - Finalizing auctions & triggering payments
 - Write volume is lower but time-critical — especially near auction end
-  Real-Time Pressure Zones:
 - Last-minute bidding: thousands of bids in final seconds
 - Real-time fan-out: pushing updates to thousands of watchers
 - Closing auctions precisely (scheduler must be accurate)
 - Payment triggers: time-sensitive + 3rd-party API dependent

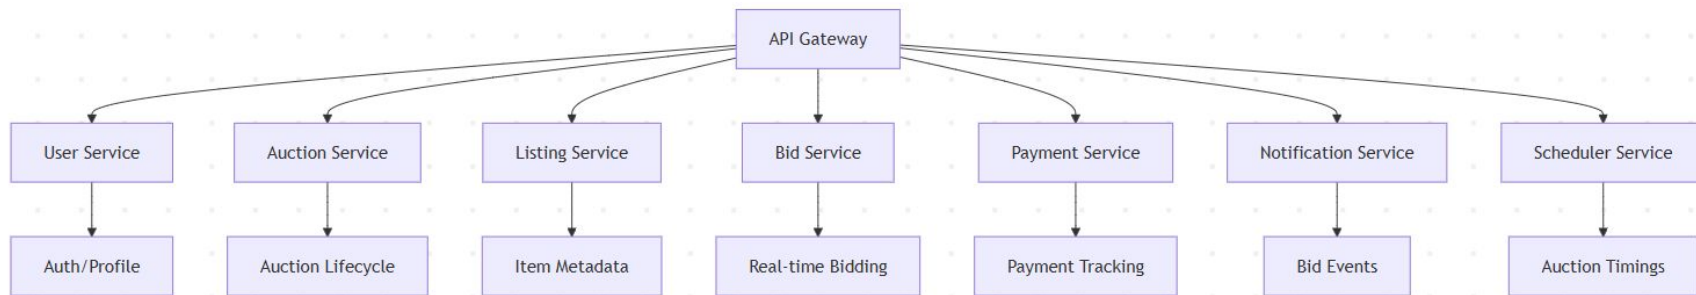
Takeaway: The system must scale for reads, but remain fast and correct for writes under pressure.

Identifying System Bottlenecks and Challenges

- Estimate per-service load to prevent surprise bottlenecks
- Isolate real-time components for better scalability
- Use async processing where consistency can be relaxed (e.g., email, payment receipt generation)
- Use horizontal scaling + partitioning for high-volume data (e.g., bids, listings)
- Plan for “hot” auctions — design for skewed load

High-Level Architecture Overview and Key components



-  API Gateway – Entry point, request routing, rate limiting
-  User Service – Auth, profile, roles (buyer/seller)
-  Auction Service – Manages auction lifecycle, bid validation
-  Listing Service – Handles item metadata, media, categories
-  Bid Service – Real-time bid management, concurrency control
-  Payment Service – Triggers post-auction payments, payment status tracking
-  Notification Service – Sends bid events, outbid alerts, auction results
-  Scheduler Service – Handles auction start/end timings
-  Analytics & Logging – Tracks usage, bids, auction trends



API Design – Key Endpoints

- User APIs:
 - POST /signup, POST /login
 - GET /user/profile
- Auction APIs:
 - POST /auctions – Create new auction
 - GET /auctions/{id} – View auction
 - POST /auctions/{id}/bids – Place bid
 - GET /auctions/{id}/bids – Bid history
 - GET /auctions/active – List active auctions
- Payment APIs:
 - POST /payments/initiate – Trigger post-auction
 - GET /payments/{id}/status
- Notification Triggers (Internal):
 - Auction ending → notify winner
 - New highest bid → notify previous top bidder

Service-to-Service Communication

-  Patterns Used:
 - Sync (REST/gRPC):
 - User auth, listing fetch
 - Auction → Bid → Payment trigger
 - Async (Pub/Sub or Event Bus):
 - New bid placed → broadcast to watchers
 - Auction ended → notify winner + trigger payment
 - Failed payment → retry/alert
-  Sample Event Topics:
 - auction.ended
 - bid.placed
 - payment.failed
 - user.registered

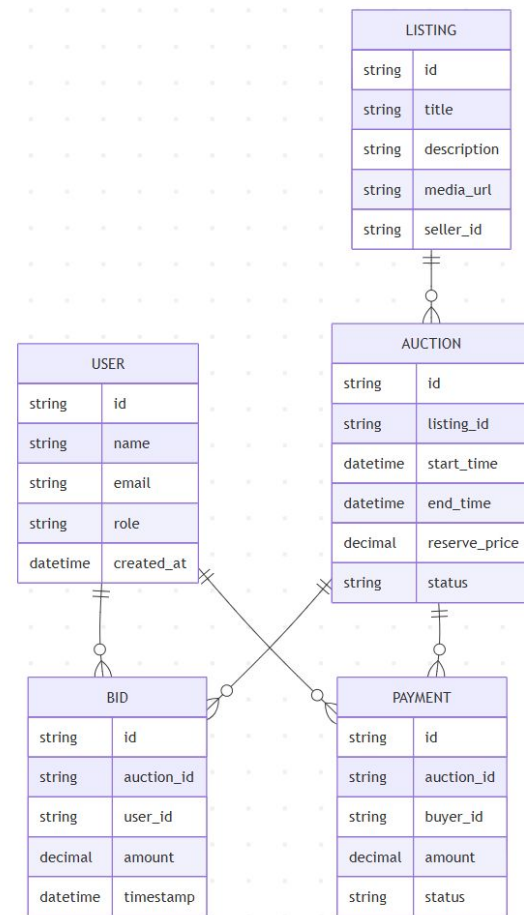
Benefits: Loosely coupled systems, retries, better failure handling

Real-Time Bid Delivery (WebSockets)

- Why WebSockets?
 - Low latency, full-duplex communication
 - Push updates to all watchers instantly
- How it works:
 - Client subscribes to auction channel: `auction:{id}`
 - Bid service validates & accepts bid → publishes event
 - WebSocket server fans out update to all connected clients
 - Key Concern: Horizontal scale under high concurrency

Data Model (Simplified)

- **USER**
 - Purpose: Stores information about the users (buyers/sellers)
- **LISTING**
 - Purpose: Contains details about the items being auctioned
- **AUCTION**
 - Purpose: Represents an auction for a specific listing, including its start and end time
- **BID**
 - Purpose: Stores bids placed on an auction by users
- **PAYMENT**
 - Purpose: Tracks payments after auction completion











Handling Auction Timers & Closures - Scheduled Job System for Auction Lifecycle

- Responsibilities:
 - Start auction at start_time
 - End auction at end_time
 - Determine winner, notify parties, trigger payment
- Design Options:
 - Dedicated Scheduler Service (cron/queue-based)
 - Use Delayed Jobs via message queue (e.g., SQS delay, Kafka, BullMQ)
 - Store timers in Redis + polling/expiration-based mechanism
- Reliability Measures:
 - Retry failed closures
 - Idempotent end-of-auction logic
 - Logging + alerts on missed timers

Strategic Tech & Infra Decisions for TinyURL

- Technology Stack for Scalability:
 - Frontend: React or Vue for responsive, real-time UI.
 - Backend: Node.js or Java with Spring Boot for handling API requests efficiently.
 - Database: PostgreSQL for relational data, Redis for caching high-demand data like auction info.
- Ensuring High Availability & Performance:
 - Load Balancing: Use cloud-based load balancers to distribute traffic evenly across servers.
 - Real-time Updates: WebSockets for real-time bidding and auction status updates.
 - Caching: Redis to cache frequently accessed data, improving performance and reducing database load.
- Security Considerations:
 - Authentication: OAuth2 for secure user login (buyers/sellers).
 - Data Encryption: SSL/TLS for encrypting data in transit, ensuring secure transactions.
- Cost Optimization:
 - Cloud Services: Use cloud providers' pay-as-you-go models to optimize infrastructure costs.
 - Autoscaling: Leverage cloud autoscaling to manage traffic surges, ensuring resource efficiency.

The Final Design - Auction Platform

-  API Gateway
-  User Service
-  Auction Service
-  Listing Service
-  Bid Service
-  Payment Service
-  Notification Service
-  Scheduler Service

