
Mastering System Design

Design a Cloud Storage System (Google Drive, Dropbox)

Understanding the Problem

- Allows users to upload, store, and manage files of any type.
- Keeps files synchronized across devices (desktop, mobile, web).
- Enables secure sharing and collaboration with others.
- Handles massive scale (millions of users, petabytes of data).
- Maintains version history and supports file recovery.



Functional Requirements

- File Upload & Download
 - Users can upload files of any type and retrieve them on-demand.
- Multi-Device Sync
 - Automatically sync files across web, mobile, and desktop clients.
- File Organization
 - Support folders, nested directories, and tagging.
- Sharing & Collaboration
 - Public/private links, editable access, shared folders.
- File Versioning
 - Maintain previous versions and allow rollback.
- Soft Delete & Restore
 - Trash bin with restore option before permanent deletion.











 *Focus on features that users expect from platforms like Google Drive or Dropbox.*

Non-Functional Requirements










- Scalability
 - Must support millions of users and petabytes of data.
- High Availability & Durability
 - 99.999999999% durability (aka “eleven nines”) for files.
- Low Latency
 - Fast uploads/downloads for small & large files.
- Security
 - End-to-end encryption, secure sharing, access control.
- Cost Efficiency
 - Optimize storage and bandwidth usage.
- Observability & Monitoring
 - Track usage, errors, sync delays, API failures.

 These qualities shape your architectural choices and trade-offs later.




Assumptions and Constraints

-  Key Assumptions (impacting architecture):
 -  Files can be large (up to 5 GB) → requires chunked uploads
 -  Users access from multiple devices → real-time sync essential
 -  Cloud object storage will be used → separation of metadata & file content
 -  Authentication & user identity is handled externally → we focus on storage layer
-  Core Constraints (that shape the system):
 -  Uploads must be resumable → chunk tracking, session management
 -  File sync latency < 5 seconds → event-driven updates, push mechanisms
 -  Storage must be cost-efficient → deduplication & lifecycle policies
 -  Fine-grained access control required → permission model and ACL support







Estimating Scale – Users, Listings, and Bookings

-  Key Metrics (Initial + Growth Projections):
 -  Users: 10M active users
 -  Average files/user: 500
 -  Total stored files: ~5B
 -  Average file size: 2MB
 -  Total data stored: ~10PB
 -  Upload rate: ~2K uploads/sec (peak)
 -  Sync events: ~10K updates/sec (peak)
-  Implication:
 - Scale requires object storage for durability, distributed metadata service, and horizontal scalability in sync and notification pipelines.








Understanding Access Patterns

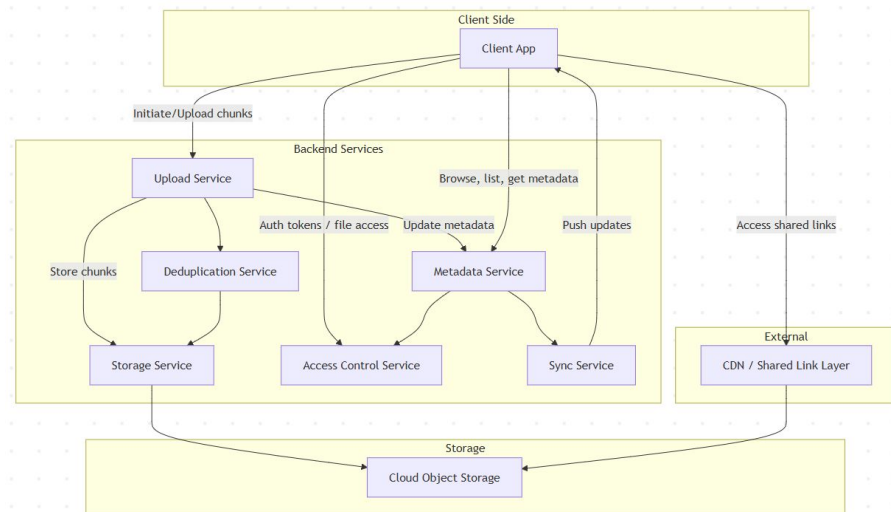
-  Write-heavy:
 - Uploads = large payloads (chunked)
 - Frequent sync updates
 - Versioning = multiple writes per file
-  Read-heavy:
 - Downloads from multiple devices
 - Folder listings & metadata reads
 - Shared link previews
-  Implication:
 - Must optimize both file write flow (chunking, session mgmt) and metadata read flow (low-latency access, caching).

Identifying System Bottlenecks and Challenges





-  Single metadata DB → becomes a hotspot
-  Large file uploads → risk of failure, timeouts
-  Real-time sync → must avoid stale states, race conditions
-  Permission checks → slow file/folder access
-  High volume shared links → unauthenticated access traffic
-  Implication:
 - Push toward partitioned metadata, chunked resumable uploads, pub-sub sync, caching for shared/public content.

Major Components in Our Cloud Storage System

-  **Upload Service:** Handles chunked/resumable uploads
-  **Metadata Service:** Stores file structure, permissions, ownership
-  **Auth Service:** Validates permissions for all operations
-  **Sync Service:** Pushes changes to connected devices in near real-time
-  **Storage Service:** Interfaces with cloud object storage
-  **Deduplication Service:** Eliminates redundant file chunks
-  **Versioning Service:** Manages previous versions and deleted files





API Design Overview

-  Upload & File Management
 - POST /upload/initiate → Starts a new file upload session and returns an Upload ID
 - PUT /upload/{id}/chunk → Uploads a chunk for the given Upload ID
 - POST /upload/{id}/complete → Finalizes upload, assembles file, triggers post-processing
-  File Retrieval & Metadata
 - GET /files/{fileId} → Downloads the file with access control checks
 - GET /files/{fileId}/metadata → Fetches file metadata, version history, permissions
-  Sharing & Collaboration
 - POST /files/{fileId}/share → Creates a shareable link (public/private)
 - GET /files/shared/{token} → Access file via shared link (unauthenticated access)
-  Sync & Change Tracking
 - GET /sync/updates → Streams or polls for real-time file/folder changes

✓ All APIs are secured, support idempotency, and follow RESTful principles. Uploads are resumable. Metadata is separate from file content.

How Services Communicate

-  Patterns Used:
 - REST + gRPC APIs for synchronous communication
 - Pub/Sub or message queues for:
 - Change events (for sync)
 - Post-upload processing
- Event sourcing for sync and versioning triggers
 -  Real-time needs → WebSockets or long polling to push file updates

Handling Chunking for Large Files

- Why Chunking?
 - Efficient upload for large files
 - Resumable uploads to prevent data loss
 - Parallel uploads for faster transfers
- How We Handle Chunking:
 - Split files into manageable chunks (e.g., 5MB per chunk)
 - Each chunk is independently uploaded (with its own checksum)
 - Chunk metadata is tracked in the Metadata Service (upload progress, chunk ID)
 - After all chunks are uploaded, they are assembled into a single file in the Storage Service
- Chunk Management:
 - Retries if a chunk fails during upload
 - Error handling: If upload is interrupted, resume from the last successful chunk

Versioning for File Updates

- Why Versioning?
 - Preserve historical file states for rollback
 - Allow users to track changes and restore previous versions
- How We Handle Versioning:
 - File metadata stores version information (timestamps, versions)
 - Each file update triggers a new version entry in the Metadata Service
 - Version IDs for easy access to specific versions
 - Users can restore any version via a simple API call (GET /files/{fileId}/version/{versionId})
- Versioning Workflow:
 - A new file upload will generate a new version if file content differs from previous
 - Change detection: Compare chunks or file hashes to detect if file has changed

Storage Strategy

- Object Storage:
 - Store files as chunks for efficient uploads, parallelism, and resilience.
 - Ensure data replication for redundancy and availability.
- SQL vs. NoSQL for Metadata:
 - SQL for structured metadata (user data, permissions, file details).
 - NoSQL for flexible or dynamic metadata (logs, dynamic file attributes).
- Scalability & Fault Tolerance:
 - Object storage offers horizontal scalability.
 - NoSQL for fast scaling of unstructured data.

Caching Strategy

- Metadata Caching:
 - Use in-memory caches for fast access to metadata.
- File Content Caching:
 - Cache files at edge locations via a CDN for quick retrieval.
- Sync & Consistency:
 - Use event systems (e.g., queues) for real-time sync and cache invalidation.

Database Schema (Rough Overview)

- User Management:
 - SQL for structured user data and relationships.
- File Metadata:
 - SQL for structured metadata, NoSQL for flexible or dynamic attributes.
- Chunk Tracking:
 - NoSQL for chunk metadata tracking.
- File Versions:
 - SQL for relational versioning, NoSQL for scalable versioning.
- Permissions:
 - SQL for managing access control (read/write).
- Audit Logging:
 - SQL for structured logs, NoSQL for high-volume logs.

Strategic Tech & Infra Decisions

- Architecture:
 - Microservices, Containers (Docker, Kubernetes)
- Storage:
 - Object Storage for file chunks (e.g., S3, MinIO)
- Database:
 - SQL (PostgreSQL) for structured metadata.
 - NoSQL (MongoDB) for scalable metadata.
- Caching:
 - In-memory caching (e.g., Redis), CDN for content delivery
- API Gateway:
 - For request routing, security (e.g., Kong, NGINX)
- Auto-Scaling:
 - Dynamic scaling based on load

The Final Design - Cloud Storage System

