




Mastering System Design

Design a Video Sharing Platform (aka YouTube)

What Are We Building?

- Design a system where users can:
 - Upload and watch videos
 - Like, comment, and share
 - Search and browse content
 - Subscribe to channels
 - View recommendations
- Key User Workflows
 -  Upload a Video
 - Chunked upload → Processing → Storage → Playback
 -  Watch a Video
 - Stream from CDN based on network/bandwidth
 -  Search/Browse
 - Filter by tags, categories, popularity
 -  Engage
 - Likes, comments, shares, subscriptions





Functional Requirements

- User registration and authentication
- Upload video content
- Encode video into multiple resolutions
- Stream videos on demand (adaptive quality)
- Video metadata (title, description, tags)
- Likes, comments, views, and subscriptions
- Search by keyword, category, or tags
- Personalized home/feed


Non-Functional Requirements

- ⚡ Low latency video streaming
- 🧩 High availability of videos and metadata
- 🚀 Scalability to billions of videos
- 💾 Efficient storage and cost management
- 🌐 Global delivery using CDN
- 🔒 Security & abuse prevention

Assumptions and Constraints

-  Assumptions
 - Users will upload mostly short-form content (≤ 15 minutes).
 - Content is static (not live).
 - Multiple video qualities will be supported (240p - 4K).
 - CDN will be used for global content delivery.
 - Metadata is small and queryable (title, tags, timestamps).
-  Constraints & Challenges
 - High storage volume: TBs per day.
 - Processing pipeline needs to scale.
 - Playback must support adaptive bitrate streaming.
 - Abuse prevention: spam, copyright, explicit content.
 - Consistency of metadata vs video availability.
 - Cost management at scale (storage + CDN egress).

Scale Assumptions

-  100M users
-  10M videos uploaded per day
-  500M daily video views
-  100M comments, likes & shares per day
-  Average video length: 10 mins
-  Video metadata: ~1KB; Engagement events: ~500B per action

Estimating Storage Needs

- Raw Video Storage
 - Avg upload: 10 min @ 5MB/min → 50MB per video
 - 10M uploads/day → 500TB/day
 - With 30-day retention (for processing): ~15PB/month
- Encoded Versions
 - Assume 4 variants: 240p, 480p, 720p, 1080p
 - Storage multiplies ~3x: 1.5PB/day → ~45PB/month






Estimating Bandwidth Needs

- Video Streaming Bandwidth
 - Avg user watches 3 videos/day → 300M hrs/day
 - Assume avg 1Mbps streaming = ~0.45GB/hr
 - Total egress/day = 135PB/day
 - CDN must handle millions of concurrent viewers
- Peak Load Estimate
 - 10M concurrent streams @ 1Mbps → ~10 Tbps egress

Metadata and Engagement Scale

- Video metadata: 10M new rows/day
- Likes/comments: 100M new events/day
- Search index updates in real time
- Hot content = read-heavy patterns


Implications of Scale Assumptions

-  **Storage**
 - Requires multi-tiered storage (hot, warm, cold)
 - Frequent writes → distributed blob storage (e.g., S3, GCS)
 - Cold storage and deletion policies to control cost
-  **Processing**
 - Encoding pipeline needs autoscaling & GPU support
 - Parallel processing jobs per resolution
-  **Global Distribution**
 - CDN integration is non-optional
 - Need region-aware content routing & geo-replication
-  **Engagement Data**
 - Write-heavy → eventual consistency and event queues
 - Aggregation for views/likes should be async and sharded
-  **Search & Discovery**
 - Near real-time indexing at scale
 - Distributed search infrastructure (e.g., Elasticsearch, Meilisearch)

Core Components Overview

- **API Gateway:** Central entry point for clients; routes requests, enforces auth, and manages rate-limiting/logging
- **Upload & Ingestion:** Handles video file uploads, generates video IDs, stores temporarily, and triggers encoding jobs via a queue
- **Encoding & Processing:** Transcodes videos to multiple resolutions, generates thumbnails, prepares HLS/DASH formats, and stores in final storage
- **Video Storage & CDN:** Manages blob storage of video chunks and manifests; integrates with CDN for fast, geo-distributed playback
- **Metadata Service:** Stores video info (title, tags, uploader, etc.), enables search/filtering, and provides mapping between video ID and location
- **User Service:** Manages user accounts, auth (OAuth/JWT), channel subscriptions, and user preferences
- **Engagement Service:** Tracks views, likes, comments; supports async event logging and anti-abuse moderation logic
- **Search & Discovery:** Powers real-time search for videos/channels using tags, titles, and trends; supports indexing of new content
- **Recommendation Engine:** Personalizes video feed using behavioral data, embeddings, and collaborative filtering models

Communication Between Services and API Design

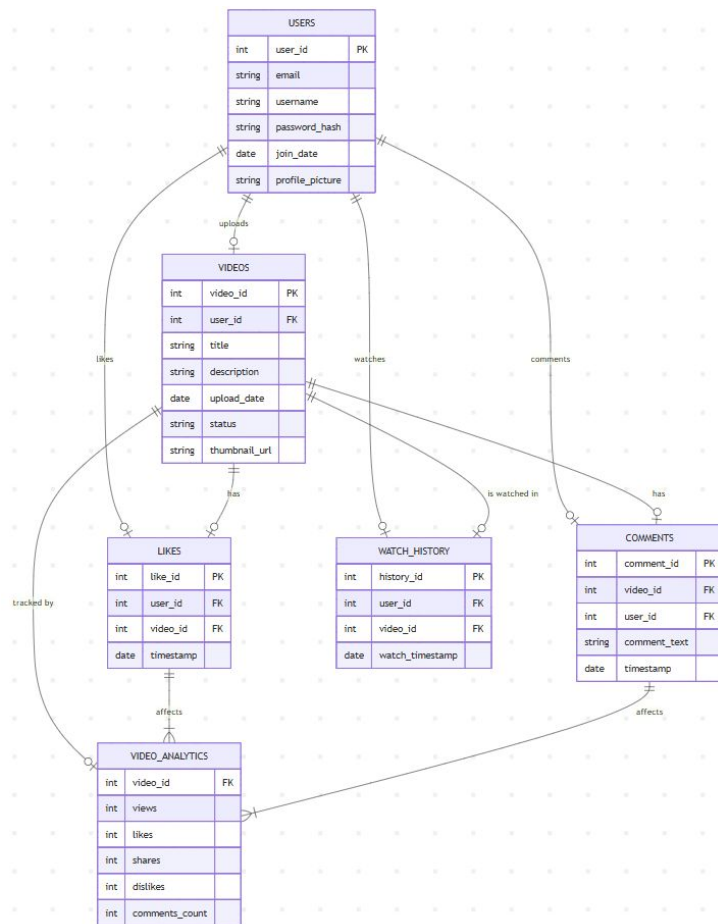
-  Types of Communication
 - Synchronous (HTTP/gRPC) for metadata fetch, user info, search
 - Asynchronous (Message Queue/Event Bus) for uploads, encoding, engagement events
- API Design
 - Client → API Gateway (REST)
 - Security: Authentication via OAuth2/JWT tokens
 - Endpoints: POST /upload, GET /videos/{id}, POST /like
 - API Gateway → Internal Services (REST)
 - Security: API Gateway handles authentication, rate-limiting, and routing
 - Services: Metadata, Encoding, Engagement, User
 - Video Upload → Encoding Service
 - Async Processing: Triggered via event bus (Kafka/SQS)
 - Security: Secure access to file storage via signed URLs

Storage and Caching Decisions

- Storage for Video Data
 - Video Files:
 - Object Storage (S3, GCS, Azure Blob) for scalable, durable storage.
 - Stores video chunks, manifest files, and thumbnails.
 - CDN:
 - Content Delivery Network (e.g., Cloudflare, Akamai) for fast, global video streaming.
 - Reduces latency by caching videos at edge locations.
- Storage for Metadata
 - Relational Database:
 - MySQL/PostgreSQL for structured metadata (video titles, tags, user data, etc.).
 - Provides fast querying and indexing for search.
 - NoSQL Database:
 - MongoDB for flexible data (e.g., user preferences, video recommendations).
- Caching & Performance
 - In-Memory Cache:
 - Redis/Memcached for frequently accessed data (video metadata, user sessions).
- Backup & Durability
 - Regular backups of databases and video files stored in geographically distributed regions.

Cursory DB Schema

- Users Table
 - user_id (PK), email, username, password_hash, join_date, profile_picture
- Videos Table
 - video_id (PK), user_id (FK), title, description, upload_date, status, thumbnail_url
- Likes Table
 - like_id (PK), user_id (FK), video_id (FK), timestamp
- Comments Table
 - comment_id (PK), video_id (FK), user_id (FK), comment_text, timestamp
- Watch History Table
 - history_id (PK), user_id (FK), video_id (FK), watch_timestamp
- Video Analytics Table
 - video_id (FK), views, likes, shares, dislikes, comments_count



Strategic Tech & Infra Decisions

- Frontend Framework
 - React.js / Vue.js for responsive, component-based UI.
- Backend Framework
 - Node.js with Express for scalable, asynchronous API handling.
- Database
 - PostgreSQL/MySQL for structured metadata storage.
 - MongoDB for flexible, scalable data (e.g., user preferences).
- Video Storage & CDN
 - AWS S3 / GCS for scalable, durable video storage.
 - Cloudflare / AWS CloudFront for low-latency video delivery.
- Authentication
 - OAuth2 / JWT for secure, token-based authentication.
- Event Processing
 - Kafka / SQS for asynchronous task processing (e.g., encoding, engagement).
- Infrastructure
 - AWS / GCP for scalable cloud infrastructure.
 - Kubernetes for efficient microservice orchestration.

The Final Design - Video Sharing Platform

