

---

---

# Mastering System Design

Design a Search Engine (aka Google)

---






---

# What Are We Building?






- Design a search engine that can:
  - Crawl and index billions of web pages
  - Serve keyword-based queries
  - Return relevant, ranked results in real-time
  - Update the index periodically for freshness
- Key Users:
  - General users performing web searches
  - Internal analytics teams querying data








# Functional Requirements

-  Web Crawling : Discover and fetch web pages at scale
-  Indexing : Extract text, normalize, and structure for fast querying
-  Keyword Search: Accept queries and return relevant documents
-  Ranking: Score and sort results by relevance
-  Re-indexing: Periodically re-crawl for updated/fresh content



# Non-Functional Requirements

-  Performance: < 200ms query response time
-  Scalability: Index billions of documents, handle 50K QPS
-  Freshness: Update index with recent changes in hours
-  Fault Tolerance: No single point of failure; retry on fetch failure
-  Storage Efficiency: Compress and deduplicate massive data volumes






# Key System Design Challenges

-  Crawling the Web: How to avoid duplication and rate limits?
-  Indexing: How to store & search through billions of documents fast?
-  Relevance: How to rank results in a meaningful way?
-  Scaling: How to scale search infra as usage and index size grow?
-  Freshness vs. Cost: How often should we re-crawl?

# Assumptions and Constraints

-  Assumptions:
  - We're focusing on publicly accessible websites (no login or dynamic JS-heavy pages).
  - Crawlers will respect robots.txt and politeness policies.
  - Query language is limited to simple keyword searches.
  - Ranking will be based on basic relevance models (TF-IDF, PageRank).
  - Users expect instantaneous results (<200ms) for top N results (e.g., 10–20).
-  Constraints:
  - Limited bandwidth for crawling — need scheduling and deduplication.
  - Storage and indexing must handle petabyte-scale data.
  - Real-time indexing not required — periodic batches acceptable.
  - No personalization or user profiling in this phase.
  - Distributed systems complexity — must handle node failures, retries, and horizontal scaling.

# Estimating Web Scale

-  Total indexed pages: ~100+ billion
-  Average page size: ~100 KB
-  Storage required: ~10 PB (raw), ~2–3 PB (compressed)
-  Daily searches: 5–10 billion (global search engines)
-  Target for MVP: 100M pages, 1M QPS capacity, 1PB storage budget

# Traffic & Query Load Estimation

- Assumptions:
  - 10M active users
  - 5 queries/user/day = 50M queries/day
  - Peak QPS = ~1000-2000 (with burst handling)
- Query Pattern:
  - 95% read (search queries)
  - 5% write (crawling/indexing updates)




# Index Size & Storage Estimation

- Raw HTML content:  $\sim 100\text{M pages} \times 100\text{KB} = \sim 10\text{ TB}$
- Processed & tokenized content:  $\sim 3\text{--}5\text{ TB}$
- Inverted index (terms  $\rightarrow$  doc IDs):  $\sim 500\text{--}800\text{ GB}$
- Forward index (doc ID  $\rightarrow$  content):  $\sim 5\text{ TB}$
- Metadata (titles, links, ranks):  $\sim 100\text{--}200\text{ GB}$

# Crawling Throughput Estimation

- Target: Crawl 100M pages in 7 days
  - Pages/day: ~14.3M
  - Pages/hour: ~600K
  - Pages/sec: ~170
- With 500 crawler workers → ~0.34 pages/sec per worker
- Feasible with politeness & backoff strategies

# Query Latency Expectations

- Query parsing: < 5ms
- Index lookup: < 20ms
- Ranking & scoring: < 50ms
- Result formatting: < 10ms
- Total SLA:  < 200ms end-to-end











# Identifying Key Bottlenecks

- **Crawling Layer:**
  - Bandwidth limitations & site rate limits
  - Risk of duplicate content
  - ➤ Solution: Distributed crawling, deduplication filters, crawl scheduling, robots.txt adherence
- **Indexing Layer:**
  - High memory & disk I/O usage when processing large datasets
  - ➤ Solution: Segment-based indexing, compression (e.g., delta encoding, front-coding)
- **Query Layer:**
  - High QPS can cause latency spikes
  - ➤ Solution: Sharded inverted indexes, result caching (popular queries), fast in-memory lookups
- **Storage Layer:**
  - Petabyte-scale storage requirement (raw + indexed data)
  - ➤ Solution: Use distributed file systems (e.g., HDFS, S3), columnar formats, cold/hot storage separation
- **Freshness Challenge:**
  - Web content changes frequently — need efficient re-crawling
  - ➤ Solution: Prioritize high-change-rate sites, adaptive re-crawling intervals, content diffing

# Scale-Driven Design Decisions

- Sharded architecture for inverted index
- Use forward & inverted indexes for speed and relevance
- Crawler partitioning by domain hash
- Replication for fault tolerance
- Introduce caching for frequent queries

# Core Components of the Search Engine

-  **Web Crawler Service**
  - Continuously fetches and downloads web pages from the internet, respecting crawl policies and rate limits.
-  **URL Frontier & Scheduler**
  - Manages the queue of URLs to be crawled, prioritizes based on freshness, domain policies, and crawl history.
-  **Content Extractor & Parser**
  - Cleans HTML, extracts visible text, metadata, and outgoing links for further processing.
-  **Indexer Service**
  - Tokenizes, stems, and builds both forward and inverted indexes from parsed content.
-  **Document Store**
  - Stores raw and parsed versions of crawled web pages along with metadata (titles, crawl date, etc.).
-  **Inverted Index Store**
  - Maps search terms (tokens) to lists of documents (docIDs) that contain them, enabling fast lookups.
-  **Query Service**
  - Parses incoming search queries, looks up inverted index, fetches candidate documents.
-  **Ranking Engine**
  - Scores candidate documents using ranking algorithms like TF-IDF, PageRank, and freshness relevance.
-  **Search API**
  - Exposes a user-facing interface (or internal API) to handle search requests and return ranked results.
-  **Cache & Frontend**
  - Stores hot query results in-memory (e.g., Redis) and provides the UI or API gateway to users.

# Crawler Coordination Architecture

## Crawler Flow:

- URL Frontier: Queue of URLs to visit (distributed, prioritized)
- Crawler Workers:
  - a. Fetch pages (respecting robots.txt, rate limits)
  - b. Store raw content in Document Store
  - c. Push URLs to fetch (links extracted) back to frontier
- Duplicate Detection: Use fingerprints (e.g., SimHash) to avoid re-crawling
- Scheduler: Controls crawl frequency and site politeness

## Coordination Strategy:

- Distributed workers pulling from sharded queues
- Hash-based partitioning by domain
- Retry & failover handling

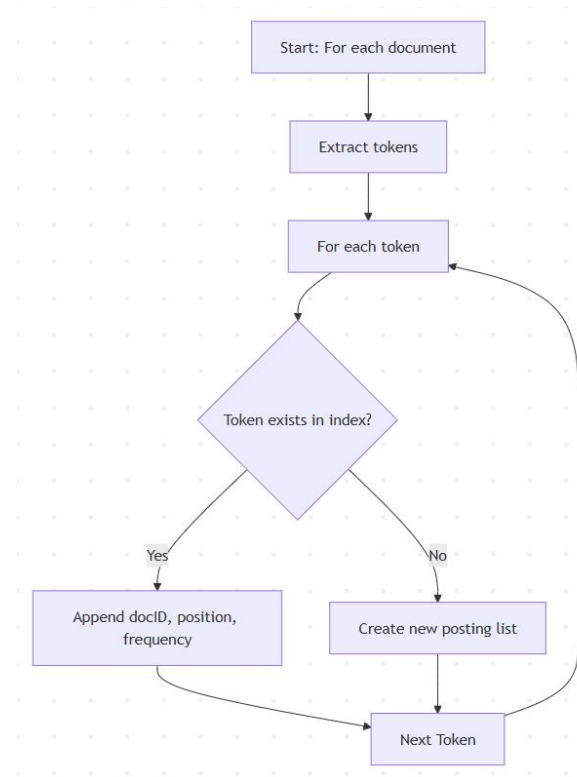


# Indexing Workflow

- Content Normalization:
  - Strip HTML, tokenize, remove stop words, apply stemming
- Forward Index Generation:
  - Stores docID → tokens + metadata
- Inverted Index Generation:
  - For each token, store list of docIDs with positions, frequencies
- Rank Feature Extraction:
  - Compute PageRank, TF-IDF, and freshness score



Indexes are partitioned and replicated across nodes






# Search Query Flow (High-Level)

[User Query] → [Search API] → [Query Parser] → [Inverted Index Lookup] → [Ranking Engine] → [Results Returned]

- Results may hit cache (for common queries)
- Top N results formatted with snippet generation



# Sample Database Structures

- Crawler DB:
  - Table: UrlQueue
- Search Index DB:
  - Table: Documents (Forward Index Table)
  - Table: InvertedIndex (Inverted Index Table)

UrlQueue		
string	url	PK
string	status	
string	domain_hash	
datetime	next_crawl_time	
int	retries	

Crawl Result

Documents		
int	doc_id	PK
string	title	
string	url	
json	tokens	
json	metadata	

Index Tokens

InvertedIndex		
string	token	PK
json	posting_list	

# Communication Between Components

- Message queues for:
  - Crawler → Parser → Indexer
  - Indexer → Index store
- HTTP/GRPC APIs for:
  - Search API → Query Service
- Caching Layer:
  - Redis or Memcached for hot queries
- Monitoring:
  - Metrics pipeline for crawl status, indexing rate, search latency

# Strategic Tech & Infra Decisions

- **Crawler Infrastructure:**
  - Distributed Crawling: Use a distributed system (e.g., Apache Kafka and Celery) to scale the crawling process and handle large volumes of URLs in parallel.
  - Cloud Providers: Leverage cloud infrastructure (AWS EC2, Google Cloud Compute) for scalability and flexibility.
- **Storage Strategy:**
  - Relational vs. NoSQL: Use NoSQL (e.g., MongoDB, DynamoDB) for the Documents and InvertedIndex tables for flexible schema and fast lookups.
  - Object Storage for Raw Data: Store raw crawled data and backups in object storage (AWS S3, GCS) for cost-effective and scalable storage.
- **Indexing Engine:**
  - Custom vs. Pre-built: Build a custom indexing engine or integrate with Elasticsearch for efficient full-text search and ranking.
  - Real-Time Indexing: Ensure the indexing process is continuous, processing documents as they're crawled for near real-time search results.
- **Ranking Algorithm:**
  - TF-IDF & PageRank: Implement TF-IDF for basic relevance and PageRank for link-based relevance scoring.
  - Freshness Consideration: Introduce freshness ranking to ensure newer content is prioritized in results.
- **API & Query Optimization:**
  - Load Balancing: Use API Gateways and Load Balancers (AWS ALB, NGINX) to distribute queries evenly.
  - Cache Results: Cache frequently searched queries using Redis or Memcached to improve performance.
- **Scalability & Fault Tolerance:**
  - Horizontal Scaling: Opt for horizontal scaling (add more servers) to handle growing data and traffic.
  - Failover & Recovery: Implement replication and automatic failover strategies to ensure high availability (using tools like Kafka or AWS RDS).

# The Final Design - Search Engine

