

---

---

# Mastering System Design

Design an E-Commerce Marketplace Platform (aka Amazon)

---

---

# What Are We Building?







- A multi-vendor e-commerce platform for buyers and sellers.
- Core functionalities:
  - Browse and search products
  - Manage inventory and product listings
  - Handle shopping cart and secure checkout
  - Process payments and detect fraud
- Must be scalable, secure, and consistent.



# Functional Requirements

- User Registration & Authentication (Buyers/Sellers)
- Product Catalog Management
- Inventory Consistency during Checkout
- Shopping Cart and Order Management
- Secure Payment Processing
- Basic Fraud Detection and Alerts
- Seller Dashboard (sales, products, payouts)
- Users:
  - Buyers: Search, browse, add to cart, pay, track orders.
  - Sellers: List products, manage stock, handle orders, track earnings.
  - Administrators: Moderate content, monitor platform health, resolve disputes.



# Non-Functional Requirements

-  **Performance:**
  - Fast product search and browsing (< 300ms latency for common queries)
-  **Scalability:**
  - Handle initial load of ~10,000 concurrent users
  - Elastic growth as user base expands
-  **Availability:**
  - 99.9% uptime target
  - Fault-tolerant core services (catalog, cart, checkout)
-  **Security:**
  - Secure user authentication (OAuth2/JWT)
  - Secure payment processing (PCI-DSS compliance best practices)
-  **Consistency:**
  - Strong consistency for inventory updates during checkout
-  **Maintainability:**
  - Modular service design (e.g., catalog service, payment service)
  - Easy onboarding of new sellers and products

# Key System Design Challenges

- Inventory Accuracy: Prevent overselling at high scale.
- High Availability: Handle millions of buyers and sellers online.
- Secure Payments: Ensure data security and prevent fraud.
- System Scalability: Grow seamlessly as user base increases.
- Efficient Search: Support quick product browsing across huge catalogs.

# Assumptions and Constraints

-  Assumptions:
  - Users have stable internet access (no offline mode required).
  - External payment gateway (e.g., Stripe, Razorpay) will handle payment processing.
  - No warehouse or logistics management initially — sellers handle shipping.
  - Product catalog size is manageable (~500K products initially).
  - Fraud detection is basic, rule-based (no ML-based advanced detection yet).
-  Constraints:
  - MVP must launch within 3–4 months.
  - Budget constraints limit infrastructure choices (e.g., cloud services preferred over heavy on-prem setups).
  - Minimal DevOps/SRE support — systems must be simple to operate.
  - Limited personalization (search is keyword-based, no ML recommendations).
  - Regulatory compliance must be ensured for payments (PCI DSS basics, GDPR for user data).

# Scale Estimations

- Users: ~1 Million registered users; ~10,000 daily active users (DAUs) at launch
- Traffic: Peak ~100 requests/second during sales
- Product Catalog: ~500,000 products
- Orders: ~1,000 orders/day initially
- Sellers: ~5,000 active sellers

# Critical System Bottlenecks

- Search and Browsing Latency: Need fast, scalable product search
- Inventory Consistency: Prevent overselling during flash sales
- Checkout and Payment Flows: Ensure reliability under peak traffic
- Database Hotspots: Popular products causing heavy reads/writes
- Fraud/Payment Errors: Detect and handle abnormal transactions quickly



# Core Services Breakdown

- **User Service** – Manage registration, login, authentication
- **Product Catalog Service** – Manage product listings, search, categories
- **Inventory Service** – Track stock levels, handle stock reservations
- **Order Management Service** – Manage cart, checkout, order lifecycle
- **Payment Service** – Handle payment processing, retries
- **Notification Service** – Send email/SMS notifications
- **Admin Service** – Seller/product moderation, platform controls

# Service APIs and Responsibilities

- User Service APIs:
  - a. POST /register, POST /login, GET /profile
- Catalog APIs:
  - a. GET /products, GET /products/{id}, POST /products
- Inventory APIs:
  - a. GET /inventory/{productId}, POST /reserve-stock
- Order APIs:
  - a. POST /cart/add, POST /checkout, GET /order/{orderId}
- Payment APIs:
  - a. POST /payments/initiate, POST /payments/verify

 APIs should be RESTful, stateless, and versioned (v1).

# Service Communication Approach

- Synchronous (HTTP REST):
  - User, Product Catalog, Inventory, Orders, Payments
- Asynchronous (Events via Message Queue):
  - Order Placed → Inventory Update → Notification
  - Payment Success → Order Confirmation → Notify Seller



✓ Use event-driven async flows for critical but non-blocking operations (e.g., stock updates, sending emails).

# Database and Storage Choices

- User Data
  - Relational DB (e.g., PostgreSQL, for strong consistency)
- Product Catalog
  - Search-optimized DB (e.g., Elasticsearch + PostgreSQL backup)
- Inventory Data
  - Relational DB (e.g., PostgreSQL, strict consistency on stock)
- Order Data
  - Relational DB (PostgreSQL, transactional)
- Payment Records
  - Relational DB (PostgreSQL, encrypted fields)
- Logs and Events
  - Object Storage (e.g., AWS S3 or similar)

- ✓ Relational databases for critical business data.
- ✓ ElasticSearch or Redisearch for fast product search indexing.








# Caching and Queueing Strategy

-  Caching (using Redis):
  - Frequently accessed product data (popular products, categories)
  - User session tokens (if server-side sessions)
  - Stock levels (read-heavy optimization)
-  Queues (using Kafka/SQS/RabbitMQ):
  - Order placement events → trigger inventory update
  - Payment success events → trigger order confirmation and notifications
  - Asynchronous notification dispatch (emails/SMS)

 Caching = Speed up reads;

 Queues = Decouple heavy async tasks from real-time flows.

# Strategic Tech & Infra Decisions

-  Database Choice:
  - PostgreSQL for core transactional data
  - Elasticsearch for product search indexing
-  Caching Layer:
  - Redis for product caching, session management, and stock-level fast reads
-  Message Queues:
  - Kafka (preferred) or AWS SQS for async event-driven flows
-  Hosting/Cloud:
  - AWS / GCP / Azure managed services (focus on managed DBs, object storage, serverless queues)
-  Authentication & Authorization:
  - OAuth 2.0 / OpenID Connect; JWT tokens for user sessions
-  Payment Integration:
  - External gateway like Stripe or Razorpay (focus on PCI compliance)
-  System Communication:
  - REST APIs for synchronous flows
  - Event-driven architecture for critical async processes

# The Final Design - ECommerce Platform

