# **Mastering System Design**

Design a Taxi Hailing App (aka Uber)

### **Understanding the Problem**

- Goal: Build a scalable taxi-hailing platform
- Core focus areas:
  - Real-time user-driver matching
  - Geo-location tracking
  - Payment processing
  - High concurrency support
- Why this is hard:
  - Requires low-latency, high-availability systems
  - o Involves mobile, backend, and external integrations





## **Functional Requirements (MVP)**

#### Rider:

- Sign up / login
- o Request a ride: source → destination
- Track driver in real time
- View ETA & trip progress
- Pay via app

#### Driver:

- Sign up / login
- Go online/offline
- Accept/reject ride requests
- Navigate to pickup & drop-off

#### System:

- Match riders with nearby available drivers
- Handle real-time location updates
- Update ride status (e.g., assigned, en route, completed)
- Calculate fares and process payments

# **Non-Functional Requirements**

- Scalability: Handle a large number of daily active users and concurrent ride sessions.
- Availability: Ensure high availability, especially during peak usage.
- Low Latency: Quick real-time matching and fast location updates.
- Data Consistency: Eventually consistent location data for real-time accuracy.
- Security: Prioritize user data privacy and secure payment processing.

#### **Assumptions and Constraints**

- Assumptions
  - Users and drivers use GPS-enabled smartphones
  - External APIs (e.g., Google Maps, Stripe) are available and reliable
  - The system launches in a single city/region initially
  - Payment is handled via in-app digital methods only
  - Real-time communication is supported via WebSockets or MQTT
- Market
   Constraints
  - o Third-party APIs (e.g., Maps, Payments) have rate limits and latency
  - Mobile connectivity is unreliable in some areas (e.g., tunnels, rural zones)
  - Drivers may go offline or change location suddenly
  - Limited compute/storage on mobile devices
  - Strict expectations for low-latency and high availability during peak usage

### Real-Time, Mapping & System Challenges

- Real-Time System Complexity
  - Match riders and drivers within <2 seconds</li>
  - Handle frequent location updates (every 2–3 seconds)
  - Maintain millions of concurrent connections reliably
- Map & Geolocation Challenges
  - Show live user/driver positions and dynamic routes
  - Perform geocoding + nearby driver search using spatial indexing
  - Deal with GPS inaccuracies and map API rate limits
- System-Level Edge Cases
  - Prevent race conditions in ride assignment
  - Handle stale/missing location data gracefully
  - Design for fallbacks when map/payment APIs fail

### **Estimating Users & Usage Metrics**

- Assume MVP launch in a single metro city.
- Estimated Daily Usage:
  - o 👩 Users: 10 million registered, 1 million DAU
  - Drivers: 200,000 active drivers/day
  - Peak concurrent sessions: ~150K (riders + drivers)
- Activity Estimates:
  - $\circ$  500,000 ride requests/day  $\rightarrow$  ~6 rides/sec
  - $\circ$  3x that for location updates  $\rightarrow$  ~18 location updates/sec
  - 1M map tile views/hour  $\rightarrow$  ~280/sec
  - 100K payment transactions/day  $\rightarrow$  ~1.2/sec

### **Key Bottlenecks & Scaling Challenges**

- Real-Time Location & Matching
  - Frequent driver updates → write-heavy, low-latency needs
  - Matching engine performance degrades as driver pool scales
  - Must push updates instantly to riders + drivers
- Mathematical Third-Party Dependency Risks
  - Map APIs: Rate limits, cost, external latency
  - o Payments: Timeouts, retries, fraud detection delays
- Platform-Wide Scaling Challenges
  - Push notifications to millions (iOS + Android)
  - Synchronizing state across services → eventual consistency
  - Cost-performance trade-offs at scale (compute, APIs, storage)

#### **Core Microservices Overview**

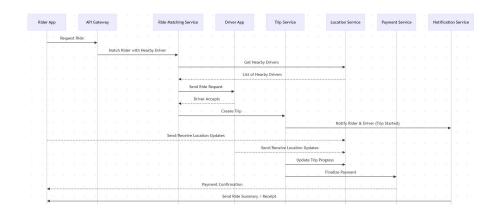
- User Service User: onboarding, auth, profile mgmt
- Driver Service: Driver onboarding, vehicle info, status
- Location Service: Real-time location ingestion & geo-indexing
- Payment Service: Ride fare calculation, payment, refunds
- Trip Management: Trip lifecycle: start, progress, end
- Motification Service: Push, SMS, and in-app alerts

### **API Gateway & Communication Patterns**

- Patterns Used:
  - Internal service-to-service: gRPC or async messaging (Kafka, NATS)
  - ⊕ External APIs exposed via API Gateway (REST/GraphQL)
  - WebSockets or MQTT for real-time location/events
- API Gateway Responsibilities:
  - Auth & rate limiting
  - Routing requests to appropriate microservices
  - Aggregating responses for frontend

#### **Client-Backend Interaction Flow**

- Rider requests a ride → via API Gateway
- 2. Ride Matching Service selects nearby driver
- 3.  $\bigcirc$  Driver gets notified  $\rightarrow$  Accepts
- 4. Cocation Service updates rider/driver positions
- 5. A Trip Service tracks ride progress
- 6. Payment Service handles post-trip payment
- 7. Notification Service sends trip updates



### **Real-Time Communication Design**

- Tech Stack:
  - Use WebSockets or MQTT for persistent rider-driver comms
  - Publish-subscribe model (e.g., Redis Pub/Sub, Kafka) for location/trip events
  - Fallback: Polling for unreliable networks
- Scenarios Handled:
  - Live driver movement on rider map
  - On Driver arriving  $\rightarrow$  pickup status  $\rightarrow$  in-progress  $\rightarrow$  completed
  - Trip cancellation, surge update, ETA changes

#### **Strategic Tech & Infra Decisions**

- Real-Time Communication
  - WebSockets vs MQTT for low-latency, persistent connections
  - o gRPC vs REST for service-to-service communication (considering latency and payload size)
- Data Storage
  - SQL (PostgreSQL/MySQL) vs NoSQL (MongoDB, Cassandra) for different use cases (e.g., user data vs ride logs)
  - Redis or Kafka for caching and event streaming
- Geospatial Indexing
  - Geohashing vs H3 for location-based queries and matching (efficiency vs accuracy)
- Scalability
  - Horizontal scaling via cloud services (AWS, GCP, Azure) vs containerization (Kubernetes) for microservices
  - Auto-scaling strategies and load balancing techniques (e.g., round-robin, least-connections)
- Fault Tolerance & High Availability
  - Multi-region replication for disaster recovery (geo-redundancy)
  - Eventual consistency vs strong consistency in trip status and payments

#### The Final Design - Taxi Hailing App

